

Causal Reasoning and Large Language Models: Opening a New Frontier for Causality

Emre Kıcıman*
Microsoft Research
emrek@microsoft.com

Robert Ness
Microsoft Research
robertness@microsoft.com

Amit Sharma
Microsoft Research
amshar@microsoft.com

Chenhao Tan
University of Chicago
chenhao@uchicago.edu

Working Paper May 9, 2023

Abstract

The causal capabilities of large language models (LLMs) is a matter of significant debate, with critical implications for the use of LLMs in societally impactful domains such as medicine, science, law, and policy. We further our understanding of LLMs and their causal implications, considering the distinctions between different types of causal reasoning tasks, as well as the entangled threats of construct and measurement validity. We find that LLM-based methods establish new state-of-the-art accuracy on multiple causal benchmarks. Algorithms based on GPT-3.5 and 4 outperform existing algorithms on a pairwise causal discovery task (97%, 13 points gain), counterfactual reasoning task (92%, 20 points gain) and actual causality (86% accuracy in determining necessary and sufficient causes in vignettes). At the same time, LLMs exhibit unpredictable failure modes and we provide some techniques to interpret their robustness.

Crucially, LLMs perform these causal tasks while relying on sources of knowledge and methods distinct from and complementary to non-LLM based approaches. Specifically, LLMs bring capabilities so far understood to be restricted to humans, such as using collected knowledge to generate causal graphs or identifying background causal context from natural language. We envision LLMs to be used alongside existing causal methods, as a proxy for human domain knowledge and to reduce human effort in setting up a causal analysis, one of the biggest impediments to the widespread adoption of causal methods. We also see existing causal methods as promising tools for LLMs to formalize, validate, and communicate their reasoning,

especially in high-stakes scenarios.

Our experiments do *not* imply that complex causal reasoning has spontaneously emerged in LLMs. However, in capturing common sense and domain knowledge about causal mechanisms and supporting translation between natural language and formal methods, LLMs open new frontiers for advancing the research, practice, and adoption of causality.

1 Introduction

Recent advances in scaling large language models (LLMs) have led to breakthroughs in AI capabilities. As language models increase in number of parameters and are trained on larger datasets, they gain complex, emergent behaviors, such as abilities to write code in programming languages, generate stories, poems, essays, and other texts, and demonstrate strong performance in certain reasoning tasks (Chen et al., 2021; Nguyen & Nadi, 2022; Bubeck et al., 2023; Katz et al., 2023; Wei et al., 2022a). Impressively, when asked to explain their outputs, update their conclusions given new evidence, and even generate counterfactuals, LLMs can create plausible responses (Nori et al., 2023; Lee et al., 2023a,b). This apparent capacity for both implicit and explicit consideration of causal factors has generated excitement towards understanding their reasoning capabilities (Hobbhahn et al., 2022; Kosoy et al., 2022; Willig et al., 2022; Liu et al., 2023; Zhang et al., 2023). Figure 2(a) shows an example of such reasoning.

At the same time, LLMs are imperfect: they can make absurd claims and are often observed to make basic errors of logic and mathematics, much less complex rea-

*Authors listed alphabetically, all contributed equally




Figure 1: When tackling real-world causal tasks, people strategically alternate between logical- and covariance-based causal reasoning as they formulate (sub-)questions, iterate, and verify their premises and implications. Now, LLMs may have the capability to automate or assist with every step of this process and seamlessly transition between covariance- and logic-based causality.

soning (Bubeck et al., 2023; Zhong et al., 2023; Ghazal et al., 2017). Figure 2(b) shows an example of an LLM making basic causal mistakes. This has incited a debate about whether LLMs truly perform causal reasoning (or not), or are simply unreliable mimics, regenerating memorized responses in the context of critical analysis and decision-making tasks across many domains, including health, medicine, legal reasoning, education, policy making and business strategy (Bender et al., 2021; Marcus, 2022). Complicating the situation is the variety of distinct formal and informal methods that people ascribe to ‘causal reasoning’, and the breadth of purposes for which people apply causal reasoning in some form.

In short, what kinds of causal reasoning might LLMs be performing, if they are, and what are the purposes for which they might be employed, if they can be? This paper clarifies this debate, advancing our understanding of LLM models and their causal implications, and proposes a framework for future research at the intersection of LLMs and causality.

We begin with a recognition of the distinct kinds of causal knowledge and reasoning implicated in this debate, including prior knowledge of general and domain-specific causal mechanisms, intuitive and logical methods for causal and counterfactual reasoning, and covariance-based causality, such as causal discovery and effect inference (see Figure 1 for an illustration). We present intriguing results in each of areas and then proceed with deeper analyses to better understand the drivers of LLM behavior and the reliability of their causal capabilities in two areas—

causal discovery and actual causality. Our results indicate that LLMs bring significant new capabilities which are complementary to existing causal methods. They do so by capturing the human domain knowledge relevant to the task, which forms an essential part of any causal analysis. As a result, LLMs have the capability of transforming how causal analysis is done, with the potential to automate or assist in each step of a causal reasoning process, shown by the arrows in Figure 1.

Summary of results. Given a set of variables, causal discovery (Peters et al., 2017) is the task of providing the causal graph characterizing their data-generating process. Starting with a pairwise causal task, we find that LLM-based methods substantially outperform existing discovery algorithms, achieving 97% accuracy on the **Tübingen benchmark** compared to the previous best of 83%. The benchmark includes over hundred causal relationships from a variety of domains, including physics, biology, zoology, cognitive science, epidemiology, and soil science. We repeat the experiment on a more specialized, real-world task related to medical pain diagnosis and obtain similarly high accuracy for GPT 3.5 and 4 models. On the more challenging task of discovering the full graph, we find that LLMs are competitive to recent deep learning-based methods. That said, LLMs do make unexpected mistakes; the surprising part is how few such mistakes are, since LLMs only consider the metadata (i.e., variable names) rather than analyzing the data values. To characterize our measurement constructs, we probe whether LLM has been trained on and memorized the datasets underlying our

SYSTEM: You are a business advisor and give clear, well-reasoned but brief (5-15 line) advice in response to questions about what people should do.

USER: A toy shop owner in the western USA wants to decide whether the ad they bought in early December is really better than their previous ads. Here is their sales data: October: \$10,200 ; November: \$10,000; December: \$13,000; January: \$10,100
They now want to decide which ad to show in February. Can you help them decide whether the increase in sales in December was due to the ads? Note that the new ad costs \$1000 more to produce, so the toy maker is interested in maximizing their profit.

GPT-4: Based on the available data, it is difficult to definitively attribute the increase in December sales to the new ad, as holiday season and other external factors may have influenced sales. To make a more informed decision, I recommend conducting an A/B test in February by dividing the targeted audience into two groups, showing one group the old ad and the other group the new ad. Compare the sales performance of both groups and calculate the net profit (considering the \$1000 additional cost for the new ad) to determine which ad is more effective in driving sales and maximizing profit. This will provide a clearer understanding of the ad's impact on sales.

(a) A GPT-4 powered advisor gives sound causal advice

USER: Thanks. In February, the toy shop owner split their audience into two groups to show them two different ads. They showed the old ad to people on the newspaper; and showed the new ad to people online. When people came into the store, the owner would ask them which ad they saw. In February, they found out that they sold \$6000 of toys to people who saw the new online ad, and \$4500 to people who saw the old newspaper ad. What should the toymaker do in March?

GPT-4: Based on the February results, the new online ad generated \$1,500 more in sales than the old newspaper ad. Considering the new ad costs \$1,000 more to produce, the net profit increase is \$500. In this case, it seems that the new online ad is more effective in driving sales. For March, the toy shop owner should continue using the new online ad, but also consider monitoring the long-term performance of the ad, as well as any potential changes in the target audience's preferences or external factors that may influence sales. Additionally, the owner may want to explore optimizing the new ad further to maximize its cost-effectiveness and reach.

(b) Continuing the conversation, GPT-4 gives a wrong causal answer

Figure 2: **Probing causal reasoning in LLMs.** Two example outputs from an LLM (GPT-4). In the first dialog, the LLM discusses causal issues, such as a potential confounder and recommends an A/B experiment to correctly characterize effects and drive the requested decision-making. The second example continues the conversation and requires arguably the same kind of causal awareness of potential confounders (e.g., the population characteristics and even population sizes of the online and newspaper audience are unknown) but the LLM proceeds regardless and provides an incorrect answer.

benchmarks, further probe the robustness of results through a wide variety of semantic and structural perturbations to our benchmark prompts, and lay out the implications for interpreting causal benchmarks.

The second major causal task we consider is actual causality (Halpern, 2016). Unlike causal discovery that deals with variables and their effect on each other, actual causality considers *individual* events and aims to find what caused them. Inferring causes of an event involves simulation of different counterfactual scenarios, but it also involves human judgment to understand the background context and determine which causes need to be considered. For the task of counterfactual reasoning, we find that LLMs like GPT-4 are capable of answering natural language questions. On a benchmark of counterfactual queries spanning basic physics, logic, and common sense, gpt-4 obtains 92% accuracy, 20 points higher than previously reported accuracy. These counterfactual capabilities also help LLMs to isolate the necessary and sufficient causes given any

event, one of the key building blocks of actual causality methods. GPT-4 obtains over 86% accuracy in identifying the necessary and sufficient causes on a commonly used benchmark of 15 vignettes, as well as a novel benchmark that avoids any memorization concerns. That said, tasks that heavily depend on an understanding of human factors in deciding the cause of an event, e.g., assessing the normality of a particular action, remain challenging for LLMs: GPT-4 obtains 70% on our benchmark task for assessing normality.

Implications for causality research. Irrespective of whether LLMs are truly performing causal reasoning or not, their empirically observed ability to perform certain causal tasks is strong enough to provide a useful augmentation for aspects of causal reasoning where we currently rely on humans alone. For example, conventional causal discovery and effect inference rely strongly on prior domain knowledge of potential causal mechanisms in a system. Current best practice is to rely on human domain experts

to provide this knowledge, yet correctly capturing domain knowledge in a formal representation suitable for analysis remains a challenge and is often a primary point of weakness for the validity of causal analyses. LLM capabilities now open the possibility of programmatic access to an array of (memorized or inferred) causal mechanisms, capturing general and domain-specific knowledge, and may augment human domain experts by aiding in bootstrapping, critiquing, etc. Other areas where LLMs provide significant benefit include the ability to understand and formalize causal scenarios, generate relevant formal premises based on background knowledge about the world; and ability to identify and correctly frame challenging causal constraints, validations and refutations, such as negative and positive controls, monotonicity relationships, etc. These are all tasks where previously we relied on human experts alone, and now can be partially or entirely automated with human supervision.

That said, LLMs do have unexpected failure modes. In each of the tasks that we studied, LLMs achieve high average accuracies but also make simple, unpredictable mistakes on certain inputs. Further, their accuracy (and consequently robustness) depends substantially on the prompt used, as observed by Long et al. (2023). We provide some basic empirical tests to probe their robustness to specific prompt language and understand the relative contribution of data memorization. However, more research is needed to understand when LLM outputs can be trusted and increase their robustness, either through external tools or other instance of LLMs themselves.

We close with an updated landscape for causal analysis research and practice. Fully characterizing LLMs inherent capacity for causal reasoning and understanding its underlying mechanisms, requires significant research efforts—and until this is accomplished, it is not prudent to trust LLMs alone in critical decision-making tasks and other causal applications. However, current capabilities are sufficiently advanced to be useful in conjunction with existing methods for formal causal reasoning, discovery, and effect inference. We discuss these opportunities, including the possibility that LLMs may provide opportunities for building tighter automated integration between logical and covariance-based approaches to causality.

Outline of the paper. Section 2 reviews background in causality and LLMs. Section 3 provides our key results on causal discovery benchmarks and the associated implications for causal discovery and causal effect inference research. Section 4 discusses actual causality and the performance of LLMs in counterfactual reasoning, determining necessary and sufficient causes, and causal judgment tasks. Finally, Section 5 discusses how these LLM capa-

bilities provide a way to augment existing causal analyses with domain knowledge and language-based reasoning, while also making it possible to combine different kinds of causal methods in a single analysis.

2 Background and Preliminaries

Working with natural language, LLMs bring capabilities to causal questions that are complementary to existing approaches. These capabilities span across hitherto disparate fields of causal enquiry, from causal discovery of general mechanisms to actual causality over events, and from statistical effect inference to logic-based reasoning. As a result, we believe that LLMs offer an opportunity to bring together different approaches to causality and build unified pipelines that seamlessly transition between language-based and data-based analyses. Below we review these different approaches or subfields under which causality is studied. We also provide an introduction to LLMs and summarize methods for probing their capabilities.

2.1 Many Kinds of Causality

TL;DR: The science of causality encompasses many ways of modeling and reasoning about cause and effect relationships. Contrasting approaches include covariance- and logic-based approaches, and methods for reasoning about type causality versus actual causality.

The science of causality is the study of cause and effect relationships, and is a fundamental tool for understanding and reasoning about the world and how it works. Correct causal reasoning is crucial to making correct decisions, building robust systems, and scientific discovery itself. While advances in causal modeling and reasoning have formalized core concepts of causality, the varying tasks and problems across domains have led different fields to use related but distinct conceptualizations and tools.

Here we describe three orthogonal categories of causal approaches and tasks. The first axis (covariance- and logic-based causality) distinguishes methodological approaches that emphasize data analysis from those that emphasize logical reasoning. The second axis (type vs. actual causality)¹ focuses on the setting for a causal question: Are we asking about causal relationships in general, such as for a population, or asking questions about specific causes for a specific event. The third axis, presented in Section 2.2

¹Type causality is sometimes referred to as general causality; and actual causality is also known as specific causality.

focus on the causal question or task itself: are we interested in discovering new relationships, characterizing the strength of specific known or hypothesized relationships, attributing blame or reward for some outcome, etc.

Covariance- and Logic-based Causality: Many fields, including statistics, biostatistics, and econometrics, place primary emphasis on *covariance-based causal analysis*. This family of methods uses statistical approaches to discover and estimate the strengths of causal relationships from data (Imbens & Rubin, 2015; Hernán & Robins, 2010). Applications include the evaluation of drug efficacy, understanding effects of novel economic policies, and optimizing business decisions. A causal analysis typically starts with a question whose answer is converted to a statistical estimand and then estimated using statistical estimators that work with available data (Pearl, 2009a).

Other domains such as law, forensic investigations, and fault diagnosis, often emphasize *logic-based causality*, which uses logical reasoning and domain knowledge to reason about causal relationships in systems (Hellner, 2000). For example, some notions of legal liability involve establishing proximate cause for an event based on reasoning about counterfactuals and plausible scenarios (Knobe & Shapiro, 2021).

Type and Actual Causality: *Type causality* encompasses inference on causal relationships between variables, such as in causal discovery and causal effect estimation (Peters et al., 2017). In contrast, *actual causality* (also called specific causality) refers to inference of the degree to which specific *events* cause other *events* (Halpern, 2016). That is, actual causality is concerned with reasoning about specific events and their causes whereas type causality focuses on variables and their average effects. For example, questions concerning medical science such as “does smoking causes lung cancer” or “what are causes of lung cancer” or “how much does smoking increase the risk of lung cancer” are examples of type-causal inference. Scientists are often interested in type causality questions, because these questions help develop theory and make predictions. For example, average causal effect of an intervention tells us something about its effect on a general population and can be used to inform policy.

In contrast, questions under the umbrella of actual causality include “Was Fred’s smoking habit responsible for his lung cancer? Or was it his exposure to asbestos?”, or “What was the reason this machine failed?”. These are questions that concern decision-making in specific situations and their answers need not generalize to other situations. Still, the answers are important to inform a specific decision or conclusion that needs to be made, e.g., deciding a legal case or fixing a machine in a factory. In

practice, beyond simple logical reasoning, to determine actual causality, people use mental “forward simulations” of processes to predict the potential outcomes of an event, or attribute the underlying events led to an observed outcome (Hegarty, 2004; Jeannerod, 2001).

2.2 Different Causal Tasks and their Connection to Kinds of Causality

TL;DR: Causal tasks, involving understanding cause-and-effect relationships, can be categorized into causal discovery, effect inference, attribution, judgment, and other tasks. Being able to solve one task does not guarantee the ability to solve others.

Causal tasks or questions seek to reason about cause-and-effect relationships in a system, often with the goal of understanding and improving a desired outcome, as other mechanisms and environment change. This broad umbrella encompasses a large array of distinct causal tasks, with its own developed methodologies. Notably, while some of these tasks intersect—they may share common abstractions (e.g., causal graphs) or the result of one task may be fed as an input of another—an ability to solve one task does not imply an ability to solve the others.

Causal discovery is the task of recovering the underlying causal mechanisms that govern a system, often described as recovering a causal graph of the system (Peters et al., 2017; Glymour et al., 2019). Causal discovery is primarily covariance-based in the context of type causality.

Effect inference is the task of characterizing the strength and shape of a known or hypothesized causal relationship. While it is most commonly characterized as relying on covariance-based methods, effect inference does rely on logical reasoning approaches for determining validity of causal assumptions (Sharma et al., 2021) and identifying approaches to validation and sensitivity analyses (e.g., negative controls (Lipsitch et al., 2010)). Effect inference is primarily focused in type causality scenarios. However, individual treatment effect estimation and counterfactual estimation straddle the realm of actual causality.

Attribution is the task of determining the cause or causes of a change. Depending on the application domain, approaches to attribution include both covariance-based and logical-based approaches, and are set in both type and actual causality settings. For example, determining the likely cause of performance slowdowns in a large-scale Internet service is usually a covariance-based analysis in a type causality setting (Budhathoki et al., 2021); determining the root cause of a specific error in a system execution can be a covariance-based analysis in an actual causality setting (Sharma et al., 2022); whereas determining the cause

of a fire in an arson investigation may rely purely on logical-reasoning in an actual causality setting (Halpern, 2016). *Judgement* tasks (Gerstenberg et al., 2014) extend attribution tasks to questions of reward or blame assignment for outcomes. They usually incorporate additional considerations of morality, normality, and intent of agents (Sloman & Lagnado, 2015).

These are only some of the many causal tasks. Others include *policy optimization*, *decision-making*, *explanation*, *scientific discovery*, and others.

2.3 LLMs and Causality

A large language model is a particular kind of machine learning model, built using transformers, a class of deep neural network architectures (Devlin et al., 2019). The primary task of LLMs is next-word completion. Initially, next-word prediction is based primarily on word distribution probabilities. Further training applies additional human feedback to shape the reward function to take into account factors beyond word distribution probability, such as instruction following and safety (Ouyang et al., 2022).

Recent work has explored causal capabilities of LLMs (Willig et al., 2022). For instance, Long et al. (2023) consider simple graphs of 3-4 nodes and test whether LLMs can recover the correct structure. They consider each variable pair (A,B) and ask an LLM to score two competing statements, one implying that A causes B and the other that B causes A. Tu et al. (2023) consider a harder task using a dataset on medical pain diagnosis and find that LLM-based discovery obtains poor accuracy. On the question of inferring causality from natural language, Hobbhahn et al. (2022) study whether LLMs can understand causal implications embedded in natural language, i.e., given two sentences, whether (the event in) one sentence is the cause of another. In this paper, we extend this line of work and make two contributions. First, we investigate graph discovery capabilities of LLMs over a broad set of complex real-world datasets and explore the robustness of LLM-based discovery. As one example, in Section 3.1, we revisit the case study on LLM-based graph discovery from Tu et al. (2023) and show that with an appropriate prompt, LLMs can achieve substantially higher accuracy (increasing the F1 score for retrieved edges from 0.21 to 0.68). Second, we probe the ability of LLMs to do counterfactual reasoning and infer necessary or sufficient cause based on natural language description of events. We also provide a unifying framework that shows how LLMs can be used to transfer knowledge between covariance-based and logic-based causal methods for a given real-world problem.

2.4 Probing LLM behaviors

TL;DR: We summarize multiple probing strategies including benchmark tests, memorization tests, and redaction tests, to evaluate the model’s performance, identify memorization issues, and analyze the model’s attention to specific aspects of input prompts.

The primary input-output mechanism exposed by an LLM is a textual prompt as input and a textual response as output. Given black-box access to an LLM, therefore, the main way to understand its causal capabilities is to probe it with different inputs and observe how its output changes. This paradigm of probing causal reasoning, however, has the classic limitation of construct validity for the measurements (Smith, 2005). That is, we may be tempted to ascribe a particular causal capability to an LLM if it answers well on a set of questions related to the capability, but the answers may not necessarily be due to the capability; they may be due to other factors such as exploiting some structure in the questions, or in the case of LLMs, memorizing similar questions that it encountered in its web-scale training set (Carlini et al., 2022). Hence we use multiple probing strategies to understand the causal capabilities of LLMs. In addition to the standard practice of testing LLM performance on existing benchmarks, we conduct memorization tests to estimate the extent to which a benchmark dataset may have been memorized, construct novel datasets to avoid memorization concerns, and conduct semantically-equivalent perturbation and other interventions on the input prompts to understand what the LLM pays attention to.

Benchmark Tests and Question-Answer Evaluation: The standard approach to evaluating LLM behavior is to ask it questions and evaluate its answers. Our question is the prompt for the LLM; and the LLM’s completion of the text is the answer. Most large scale benchmarks and evaluations follow this approach, using purposefully created benchmarks and/or adapting standardized exams written to assess human capability (Ghazal et al., 2017; Nori et al., 2023; Zhong et al., 2023).

Our evaluation approach to characterizing the causal capabilities of LLMs begins with such question-answering evaluations. For each evaluation, we ask a series of questions or causal challenges and score the correctness of the resulting answer. Then, we probe deeper, to better understand the threats to the validity of the question-answer evaluations, and the likely robustness of LLM capabilities. **Memorization Test²:** The primary threat to the validity of benchmark or question-answer style assessments is that

²This method for testing for memorization and prior exposure of the

the LLM may have directly memorized the benchmark answers. In this case, our questions are likely not testing the LLM’s inherent capability to complete a task (unless memorization is the capability we are testing!)

To test whether the LLM has memorized a particular dataset or benchmark, we give the LLM a partial row of data from the dataset and ask it to autocomplete the remainder of the row. For a question-answer benchmark, we give the LLM half of the question, and ask it to auto-complete the remainder of the question itself. To encourage the LLM to succeed, we prepend details about the dataset, such as its name, URL, and description, and also provide few-shot examples. The final measurement from the memorization test is the percentage of rows the LLM was able to regenerate correctly. Supplementary E.1 provides additional details on the procedure.

Redaction Test: When we see an LLM generate a correct answer to a question, we are still sometimes unsure why. Is the LLM attending to the appropriate aspects of the question, or has it learned to repeat unreliable patterns that may lead to errors in the future. To better understand what aspects of the prompt or question an LLM is attending to, we use redaction and perturbation tests, motivated from explainable AI methods for NLP Sinha et al. (2021); Danilevsky et al. (2020). First, we redact words in the prompt, one by one, and see how the answer changes each time. Changes in the answer indicate the LLM is attending to the redacted word.

3 LLMs and Causal Discovery

TL;DR: LLMs can achieve competitive performance in determining pairwise causal relationships with accuracies up to 97% but their performance varies depending on prompt engineering and may occasionally be inconsistent. Extending knowledge-based causal discovery to full graph discovery poses additional challenges, but can also achieve strong results.

Causal discovery is the task of learning the structure of causal relationships among variables Peters et al. (2017). Typically, the output of causal discovery is a directed graph where the edges denote the presence and direction of causal effect. Such a graph characterizes the underlying data-generating process (DGP) for a dataset and specifies how a change in one variable may (or may not) affect the others. This graph is then used as a base on which downstream analysis relevant to a task is conducted, such as for effect

inference, prediction or attribution. Having the correct graph that encodes causal assumptions is critical for ensuring the correctness of any downstream analysis. In other words, causal discovery is a fundamental task that every other causal task depends on.

The challenge, however, is that it is generally not possible to learn the correct graph for a given dataset, given only observational data. The reason is that multiple graph structures are equally likely given the same data distribution, a set of graphs known as the Markov equivalence class (Pearl, 2009c). In the last two decades, two main approaches have been proposed to overcome this limitation. The first is to restrict data-generating process to specific functional forms under which identification of a single graph is possible (Glymour et al., 2019). In some specific settings, such as adding non-gaussian noise to linear data-generating process (Shimizu et al., 2006) or assuming that all functions are non-linear with additive noise (Zhang & Chan, 2006; Zhang & Hyvärinen, 2009), identification is possible. However, there still exist simple settings that are unidentifiable, e.g., a dataset with linear equations and gaussian noise. The second approach is to utilize the power of deep learning to model covariances of all variables jointly and hope that it improves the quality of learnt graphs. However, the identification issue is still not resolved. As a result, recent evaluations of state-of-the-art causal discovery methods on real-world datasets present a sobering picture of their effectiveness (Kaiser & Sipos, 2022; Tu et al., 2019; Huang et al., 2021).

LLMs offer a fresh perspective on the causal discovery problem by focusing on the *metadata* associated with variables in a dataset, rather than their data values. Typically, such metadata-based reasoning is done by human domain experts when they construct causal graphs. By looking at the names of variables and the problem context, for example, people can construct a causal structure based on their knowledge of physics, common sense, or specialized domain knowledge. However, this is a challenging process, hindering the widespread use of such metadata-based reasoning.

We find that LLMs can bridge this gap by filling in the domain knowledge that earlier only humans could provide. In contrast to causal discovery algorithms that use data values of variables, LLMs can infer causal structure by reasoning on *metadata* associated with the variables, for example, the name of the variable and the problem context expressed in natural language. In other words, similar to how domain experts formalize their knowledge in a graph, LLMs can use the knowledge from their training data to infer the edges of the causal graph. To differentiate from the existing covariance-based causal discovery, we call the

LLM to data was developed by Harsha Nori and Rich Caruana, and will be presented in more detail in a soon-to-be released paper.

Variable A	Variable B	Domain
Age of Abalone	Shell weight	Zoology
Cement	Compressive strength of concrete	Engineering
Alcohol	Mean corpuscular volume	Biology
Organic carbon in soil	Clay content in soil	Pedology
PPFD (Photosynthetic Photon Flux Density)	Net Ecosystem productivity	Physics
Drinking water access	Infant mortality	Epidemiology
Ozone concentration	Radiation	Atmospheric Science
Contrast of tilted Gabor patches	Accuracy of detection by participants	Cognitive Science
Time for 1/6 rotation of a Stirling engine	Heat bath temperature	Engineering
Time for passing first segment of a ball track	Time for passing second segment	Basic Physics

Table 1: Example cause-effect pairs from the Tübingen benchmark. The task is to determine whether Variable A causes Variable B, or vice-versa.

LLMs capability as *knowledge-based* causal discovery.

Remarkably, LLMs practicing knowledge-based discovery outperform state-of-the-art covariance-based algorithms on causal discovery benchmarks. Below we present experiments on inferring causal structure using LLMs. We first start with pairwise causal discovery: a simple task involving two variables where the goal is to decide the direction of causal effect, across a range of domains. Next, we will study the problem of full graph discovery in two datasets: one in medicine and one in climate science. While the first task includes many “simple” relationships that an average person is expected to answer correctly, the other two tasks require specialized knowledge (in neuro-pathic pain and arctic atmosphere science respectively) that a lay person would not be able to infer.

3.1 Pairwise causal discovery: Inferring causal direction among variable pairs

TL;DR: LLMs enable knowledge-based causal discovery, and achieve competitive performance in determining pairwise causal relationships between variables, across datasets from multiple domains, including medicine and climate science.

We start with the pairwise causal relationship discovery task (Hoyer et al., 2008). Typically, these tasks are set up as discovering causal relations between two variables based on observed data. In this work, instead of relying on observed data, we directly ask large language models whether a variable causes another variable.

3.1.1 Tübingen cause-effect pairs dataset

This dataset (Mooij et al., 2016) consists of data for 108 different cause-effect pairs selected from 37 datasets from various domains, including meteorology, biology, medicine,

engineering, and economics (see examples in Table 1). As directionality of an edge is a fundamental building block for learning the full graph, it is a widely used dataset for benchmarking causal discovery algorithms. However, the Markov equivalence class of graphs admissible for a two-node graph (A, B) contains both $A \rightarrow B$ and $B \rightarrow A$. Therefore, the dataset remains a challenging one for causal discovery algorithms, with many recent methods achieving less than 80% accuracy, as shown in Table 2. The best known accuracy is 83%, achieved by the Mosaic algorithm (Wu & Fukumizu, 2020).

We now apply LLMs to the pairwise discovery problem. We extract the names of each variable from the benchmark and use them to construct prompts for the LLMs. Example prompts are shown in Suppl. A.1 (Table 14). We start with a simple prompt that asks, “Does changing A cause a change in B?”, where A and B are the variable names. For a given pair (A, B) , we ask the question in both directions and then take the mean accuracy. We choose the mean accuracy since it maps correctly to the accuracy of any method choosing between $A \rightarrow B$ or $B \rightarrow A$. Specifically, the LLM-based method obtains an accuracy of 0 if it answers both questions incorrectly. If it answers one of the questions correctly (indicating that none cause each other or both cause each other), then it effectively chooses the output at random ($A \rightarrow B$ or $B \rightarrow A$). We also report weighted accuracy, as recommended by Mooij et al. (2016) to avoid the effect of overcounting some similar pairs.

Results. The second half of Table 2 shows the performance of different LLMs on the task. Similar to prior studies on capabilities of LLMs (Wei et al., 2022a), we observe the emergent behavior that only text-davinci-002 and above achieve a non-trivial accuracy than random chance. This suggests that the ability of inferring causal direction emerges as the size of the model is increased. With GPT3.5

class models (text-davinci-003 and gpt-3.5-turbo), accuracy of LLM-based method reaches 83% and is competitive to the covariance-based causal discovery algorithms.

Following recent work showing the value of the prompt tuning for obtaining high-quality results (Wei et al., 2022b; Long et al., 2023), we employ two different kinds of prompts. In the first, we prepend the prompt with the message, "You are a helpful assistant for causal reasoning", with the intention of steering the LLM's output space towards more causally consistent answers. This simple message of *asking* gpt-turbo-3.5 (i.e., ChatGPT) to be an agent of causal reasoning increases the (weighted) accuracy by almost 5% to 86.9%, consistent with the impact of prompt engineering observed by (Long et al., 2023). The accuracy is higher than previously reported LLM-based accuracy on the task from Choi et al. (2022) that used the davinci-instruct-beta model (83% unweighted accuracy). Moreover, by changing the prompt to instead ask a single question that asks the LLM to output the more likely causal direction between $A \rightarrow B$ or $A \leftarrow B$ while explaining its reasoning in a "step-by-step" manner (see Suppl. A.1), the weighted accuracy improves to 92.1%. Finally, using gpt-4 with this prompt, the weighted accuracy shoots up to 97%. This accuracy is substantially higher than that of the best-known covariance-based causal discovery method.

Understanding LLM output. The combined single prompt allows us to inspect the LLM output, since it provides an explanation in addition to its final answer. For the examples below, we show the output from gpt-3.5-turbo. Let us look at the example, (age of abalone, length of abalone) where the LLM correctly outputs that changing the age of an abalone is more likely to change its length. Figure 3(a) shows the LLM's output. Its explanation gives the impression that it can *reason* about the problem using its knowledge to arrive at the correct answer:

However, on another pair that arguably should involve similar reasoning, (age of abalone, diameter of abalone), the LLM outputs an incorrect answer and its explanation does not make much sense (Figure 3(b)).

The fact that the LLM can make such errors is not surprising, since it is not explicitly optimizing for being consistent with prior knowledge or causal mechanisms. However, it is remarkable that such errors occur infrequently: The LLM made only 12 errors in the dataset of 108 pairs. On some other pairs where the LLM's output did not match the ground-truth, it reveals possible ambiguity in the variable names, that can be fixed by adding more context to the question. Consider the example of (ozone concentration, radiation) below where the LLM initially provides an incorrect answer.

LLM: 1. Ozone is a gas in Earth's atmosphere that absorbs

Model	Acc.	Wt. Acc.
Slope (Marx & Vreeken, 2017)	0.75	0.83
bQCD (Tagasovska et al., 2020)	0.68	0.75
PNL-MLP (Zhang & Hyvarinen, 2012)	0.75	0.73
Mosaic (Wu & Fukumizu, 2020)	83.3	81.5
ada	0.50	0.50
text-ada-001	0.49	0.50
babbage	0.51	0.50
text-babbage-001	0.50	0.50
curie	0.51	0.52
text-curie-001	0.50	0.50
davinci	0.48	0.47
text-davinci-001	0.50	0.50
text-davinci-002	0.79	0.79
text-davinci-003	0.82	0.83
LMPrior (Choi et al., 2022)	0.83	-
gpt-3.5-turbo	0.81	0.83
gpt-3.5-turbo (causal agent)	0.86	0.87
gpt-3.5-turbo (single prompt)	0.89	0.92
gpt-4 (single prompt)	0.96	0.97

Table 2: Accuracy of different versions of GPT on the Tübingen cause-effect pairs dataset. The best LLM performance outperforms the current state-of-the-art covariance-based approaches that rely on observational data of the two variables. Weighted accuracy weights individual pairs to account for overcounting due to some pairs sharing the same source dataset. The causal agent is gpt-3.5-turbo with system message set as "You are a helpful assistant for causal reasoning.". LMPrior uses davinci-instruct-beta.

most of the Sun's ultraviolet (UV) radiation.

2. Chlorofluorocarbons (CFCs), pollutants that were once widely used in refrigerants and aerosol sprays, have been found to destroy ozone molecules in the atmosphere.

3. The destruction of the ozone layer in the upper atmosphere causes more UV radiation to reach the surface of Earth.

4. Increased UV radiation can lead to various negative effects such as skin cancer, cataracts, and damage to crops.

From these facts, we can determine that a decrease in ozone concentration would lead to an increase in radiation. Therefore, Option A is more likely:

<Answer>A</Answer>

The correct answer according to the benchmark is that radiation causes ozone concentration. However, the LLM interpreted the question as about the ozone concentration in the stratospheric ozone layer, whereas the question was about ground-level ozone concentration measured in a city. After adding this context, the LLM's answer changes,

SYSTEM: You are a helpful assistant for causal reasoning.

USER: Which cause-and-effect relationship is more likely?

A. changing the age of an abalone causes a change in its length.

B. changing the length of an abalone causes a change in its age.

Let's work this out in a step by step way to be sure that we have the right answer. Then provide your final answer within the tags <Answer>A/B</Answer>.

GPT-3.5-TURBO: 1. Cause: Changing the age of an abalone
Effect: a change in its length.

2. Cause: Changing the length of an abalone
Effect: a change in its age.

Based on the biological knowledge about abalones, it is more likely that changing the age of an abalone causes a change in its length. This is because the length of an abalone is known to increase as it ages, and not the other way around. Therefore, the correct answer is <Answer>A</Answer>.

(a) Correct answer

SYSTEM: You are a helpful assistant for causal reasoning.

USER: Which cause-and-effect relationship is more likely?

A. changing the age of an abalone causes a change in its diameter.

B. changing the diameter of an abalone causes a change in its age.

Let's work this out in a step by step way to be sure that we have the right answer. Then provide your final answer within the tags <Answer>A/B</Answer>.

GPT-3.5-TURBO: Step 1: Determine which variable is the cause and which is the effect.

In Option A, "age" is the cause and "diameter" is the effect.

In Option B, "diameter" is the cause and "age" is the effect.

Step 2: Consider which variable is more likely to cause a change in the other.

Abalone growth is primarily determined by environmental factors such as water temperature, food availability, and water quality. It is unlikely that changing the age of an abalone would cause its diameter to change significantly. On the other hand, changes in the environment could potentially impact the size of the abalone, and therefore, its diameter could change. Therefore, Option B is more likely: changing the diameter of an abalone causes a change in its age.

<Answer>B</Answer>

(b) Incorrect answer

Figure 3: **Probing reasoning of LLMs for deciding causal direction.** Two example outputs from a LLM (gpt-3.5-turbo). The first gives the impression of the LLM performing causal reasoning to arrive at the correct answer—that age may cause the length of abalone to increase. The second example asks a question that requires arguably the same causal knowledge but the LLM proceeds with an incoherent argument and provides an incorrect answer.

USER: Wait, this question is not about stratospheric ozone concentration. Rather, I'm asking about the ground-level ozone concentration measured in a city and the corresponding ground-level radiation due to sunlight.

GPT-3.5-TURBO: I see, thank you for clarifying. In that case, the answer is likely option B: changing radiation causes a change in ground-level ozone concentration. This is because ground-level ozone is formed through a chemical reaction between nitrogen oxides (NOx) and volatile organic compounds (VOCs) in the presence of sunlight. This reaction is known as photochemical smog and occurs more frequently on hot, sunny days when there is more radiation. Therefore, an increase in radiation would likely lead to an increase in ground-level ozone concentration.

<Answer>B</Answer>

Finally, in some cases, the direction of the cause-effect relationship is not absolute and depends on other supporting information. For the pair (stock return of Hang Seng bank, stock return of HSBC Hldgs.), the LLM correctly determines that Hang Seng is a subsidiary of HSBC Hldgs. but

its answer does not match the benchmark's label. HSBC Hldgs. owns a majority stake in Hang Seng bank, so the benchmark answer indicates that any change in Hang Seng's stock return should also affect HSBC Hldgs. stock return. At the same time, since stock markets do not follow any fixed pattern, it is possible that a major change in HSBC Hldgs. stock return causes people to invest more or less in Hang Seng Bank and hence affect its stock return, as the LLM output argues.

LLM: To determine which cause-and-effect relationship is more likely, we need to consider the nature of the two companies and their relationship with each other.

Hang Seng Bank and HSBC Holdings are two separate companies but they are both major players in the Hong Kong financial market. Hang Seng Bank is a subsidiary of HSBC Holdings and is listed on the Hong Kong Stock Exchange. HSBC Holdings is a multinational banking and financial services company that operates in over 60 countries worldwide.

Given that relationship and the fact that Hang Seng Bank is a subsidiary of HSBC Holdings, it is more

Variable A	Variable B	Dir.
Right L1 Radiculopathy	Right adductor tendonitis	→
Pharyngeal discomfort	Right C3 Radiculopathy	←
Right L5 Radiculopathy	Lumbago	→
Left PTA	Left L4 Radiculopathy	←
Left T3 Radiculopathy	Toracal dysfunction	→
DLS L5-S1	Right S1 Radiculopathy	→
Left C3 Radiculopathy	DLS C2-C3	←
Left C7 Radiculopathy	Left medial elbow problem	→
Right Ischias	Right L5 Radiculopathy	←
Right Morton trouble	Right L5 Radiculopathy	←

Table 3: Example cause-effect pairs from the Neuropathic pain diagnosis benchmark. ‘Dir.’ refers to the ground-truth causal direction between the variables.

likely that:

<Answer>B</Answer>changing the stock return of HSBC Hldgs causes a change in the stock return of Hang Seng Bank.

3.1.2 Neuropathic pain dataset

We now consider a ground-truth causal graph constructed by medical experts on neuropathic pain diagnosis (Tu et al., 2019). It contains relationships between different nerves and the associated symptoms that patients express. Table 3 shows example of edges in the graph. Unlike the Tübingen benchmark which contains some commonly known variables, understanding the names of variables in this graph requires specialized medical knowledge.

We evaluate LLMs on a pairwise causality task similar to the Tübingen benchmark. We do so by creating a pair for each edge in the graph and setup the task to determine the direction of the edge. That is, for all the edges present in the ground-truth graph, we ask an LLM to infer the correct direction of causality. We obtain a dataset of 475 edges.

Table 4 shows the accuracy of different LLMs in determining the direction of edges. Similar to the Tübingen benchmark, smaller language models are unable to obtain accuracy substantially higher than random chance (50%). In addition, even text-davinci-003, the largest text completion model in our comparison set, is unable to obtain a high accuracy using the two-sided prompt. However, chat-based models, trained using human feedback, are able to distinguish causal direction accurately. With the two-sided prompt, gpt-3.5-turbo achieves the highest accuracy of 75%. The choice of prompt also has a big impact. As with the Tübingen dataset, using single prompt per variable pair leads to a significant increase in accuracy: text-davinci-003

Model	Accuracy
ada	40.1
text-ada-001	50.0
babbage	50.0
text-babbage-001	50.9
curie	50.0
text-curie-001	50.0
davinci	38.4
text-davinci-001	50.0
text-davinci-002	51.7
text-davinci-003	55.1
gpt-3.5-turbo	71.1
gpt-3.5-turbo (neuropathic pain expert)	75.1
gp4-4	78.4
gpt-4 (neuropathic pain expert)	84.3
text-davinci-003 (single prompt)	86.0
gpt-3.5-turbo (single prompt)	85.5
gpt-4 (single prompt)	96.2

Table 4: Accuracy of different versions of GPT on the inferring the edge directions of the Neuropathic pain diagnosis graph. As with the Tübingen dataset, LLMs like gpt-3.5-turbo obtain more than 85% accuracy on determining the direction of edges. The causal agent is gpt-3.5-turbo with a system message set as “You are a helpful assistant for causal reasoning.”

now matches gpt-3.5-turbo and both models obtain more than 85% accuracy. While the performance of these models is impressive, gpt-4 obtains an even higher accuracy (96%), yielding more than 457 correct responses out of 475.

Understanding LLM output. We use the output of gpt-3.5-turbo (single prompt) to understand how LLMs obtain such high accuracies. On manual inspection of the output by gpt-3.5-turbo, we find that it correctly understand the medical terms in almost all cases. In the example below, the LLM provides the correct definitions of DLS T5-T6 and Left T6 Radiculopathy (without access to any such information in the prompt) and provides the correct causal direction. Suppl. A.2 provides the LLM output for this pair. However, even as it provides the correct definitions, it can provide erroneous output in unpredictable ways. Here it went from the correct definition of the same concept (DLS) to confusing it with the “T5-T6 region” associated with DLS T5-T6 (see Suppl. A.2 for the full output).

Moreover, the example in Table 5 adds evidence that LLMs may not be capable of consistent and coherent reasoning, even as they output correct answers for a majority of the causal questions. In the left panel, the LLM provides the correct reasoning for most of its answer but ultimately chooses a final answer inconsistent with its own reason-

ing (A instead of B). Unlike reasoning errors, fortunately self-consistency is an easier problem to verify using the LLM itself. We find that gpt3.5-turbo and earlier class of models are unable to check consistency of an output, but GPT-4 is able to. Therefore, one can pass all LLM outputs through a GPT-4 verifier to ensure self-consistency of the outputs. The right panel of Table 5 shows the verifier prompt and output from GPT-4 (the question and answer from gpt-3.5-turbo are not fully shown for brevity).

In addition to self-consistency, when LLM outputs do not match the ground-truth, they can sometimes help critique the ground-truth itself. For instance, the ground-truth graph contains an edge from Left L5 Radiculopathy to Obesity. However, the LLM output provides the opposite direction: obesity can cause L5 Radiculopathy, which is corroborated by research that claims obesity as a risk factor for radiculopathy (Atchison & Vincent, 2012; Tomić et al., 2009).

LLM: Step 1: Understand the terms

- Neuropathic pain: pain caused by damage or disease affecting the nervous system
- L5 Radiculopathy: a condition where the nerve root that exits the spine at the L5 level is compressed or irritated, causing pain, numbness, or weakness in the leg
- Obesity: a medical condition characterized by excessive body fat that increases the risk of health problems

Step 2: Analyze the cause-and-effect relationship

- Option A: Left L5 Radiculopathy causes Left Obesity. This relationship does not make sense because L5 radiculopathy is a nerve condition that affects the leg, while obesity is a whole-body condition related to excessive body fat. It is unlikely that a leg nerve condition would cause obesity.
- Option B: Left Obesity causes Left L5 Radiculopathy. This relationship is more plausible because obesity can put extra pressure on the spine and nerves, leading to nerve compression or irritation. It is possible that obesity could cause L5 radiculopathy.

Step 3: Choose the more likely cause-and-effect relationship

- Based on the analysis, Option B is more likely: Left Obesity causes Left L5 Radiculopathy.

<Answer>B</Answer>

3.1.3 Other properties of pair-wise relationships

In addition to asking the LLM about the existence and direction of a causal relationship between a pair of edges, we can also ask many other questions as well. For example, we can ask if a relationship may be confounded by other variables, or what variables may mediate or moderate the relationship. We can ask whether a causal relationship is

monotonic, whether its effects are homogeneous or not with respect to other properties. Other useful properties to identify include the time-lag of a causal relationship, whether it is stable, and whether it is subject to spillover effects.

3.2 Full graph discovery

TL;DR: Extending knowledge-based causal discovery to full graph discovery poses additional challenges, such as distinguishing between direct and indirect causes. With the right prompt, LLMs provide non-trivial utility for inferring full causal graphs.

We now extend the pairwise analysis to discovering the full graph. Given a set of variables, a straightforward extension is to construct a list of all possible variable pairs and repeat the pairwise test for each pair. However, full graph discovery provides additional challenges. First, the pairwise analysis assumed the existence of an edge and our goal was simply to determine its direction. In the graph discovery problem, there are three possible options for each variable pair: $A \rightarrow B$, $A \leftarrow B$, or no edge exists. Second, graph discovery requires the ability to distinguish between direct and indirect causes, given the other nodes in the graph. For example, if the true relationship is $A \rightarrow B \rightarrow C$, then it may be correct to output both $A \rightarrow B$ and $A \rightarrow C$ in a pairwise task, but for graph discovery, outputting $A \rightarrow C$ is an incorrect answer. Moreover, the decision depends on which variables are present in the input variable set for the graph discovery task. For example, if B is not one of the variables in the input variable set, then $A \rightarrow C$ is a valid edge and should be included.

Due to these reasons, it is unclear how well a simple extension of the pairwise task using LLMs will transfer to the full graph discovery problem. For simple graphs over 3-4 variables, Long et al. (2023) show that LLMs can obtain promising accuracy in inferring the edges. Below we report the results of applying the pairwise LLM prompts on more complex, real-world datasets involving specialized knowledge: on pain diagnosis and on arctic sea ice coverage.

3.2.1 Neuropathic pain dataset

We continue our exploration with the neuropathic pain dataset. It has 221 variables, so the total number of possible variable pairs are $C(221, 2) = 24310$. Before conducting full graph discovery, we utilize a 100-pair dataset provided by the authors of the dataset to evaluate LLM’s capability to discover causal relationships (Tu et al., 2023). It has 50 variable pairs that form edges in the graph and 50 pairs that

Initial response	Critiquing the response
<p>SYSTEM: You are an expert on neuropathic pain diagnosis.</p> <p>USER: You are a helpful assistant to a neuropathic pain diagnosis expert. Which cause-and-effect relationship is more likely?</p> <p>A. Right C7 Radiculopathy causes DLS C6-C7.</p> <p>B. DLS C6-C7 causes Right C7 Radiculopathy.</p> <p>Let’s work this out in a step by step way to be sure that we have the right answer. Then provide your final answer within the tags <Answer>A/B</Answer>.</p> <p>GPT-3.5-TURBO: <i>First, we need to understand what each term means. DLS stands for degenerative lumbar spondylosis, which is a condition where the spinal discs in the lower back degenerate and cause pain. C6-C7 refers to the sixth and seventh vertebrae in the cervical (neck) spine. Radiculopathy is a condition where a nerve root is compressed or damaged, causing pain, numbness, or weakness in the corresponding area of the body.</i></p> <p><i>Based on this information, it is more likely that option A is the correct cause-and-effect relationship. This is because radiculopathy is often caused by compression or damage to a nerve root, which could be the result of degenerative changes in the spine (such as DLS). In this case, the degeneration in the lower back (DLS) could be causing compression or damage to the nerve root at the C7 level, resulting in right C7 radiculopathy.</i></p> <p><i>Therefore, the answer is <Answer>A</Answer>.</i></p>	<p>USER: Analyze the output from an AI assistant. Is the final answer (A) consistent with the reasoning provided by the assistant?</p> <p>Question: { {USERPROMPT-FROM-BEFORE} }</p> <p>AI assistant: { {ANSWER-FROM-BEFORE} }</p> <p>GPT-4-WEB: <i>The final answer (A) is not consistent with the reasoning provided by the assistant. The reasoning suggests that DLS C6-C7 could cause Right C7 Radiculopathy, which corresponds to option B. The correct answer should be <Answer>B</Answer>.</i></p>

Table 5: **Critiquing LLM output using another LLM instance.** To increase robustness of an LLM’s response, we can use GPT-4 as a critic. The left panel shows an incorrect reply from gpt-3.5-turbo wherein the reasoning is correct but the LLM outputs the incorrect option (A). We create a special “critique” prompt that asks gpt-4 to evaluate the response from an AI assistant for self-consistency. gpt-4 finds the logical inconsistency and provides the correct answer.

Model	Precision	Recall	F1
Random	0.25	0.5	0.33
chatGPT (Tu et al., 2023)	1	0.12	0.21
text-davinci-003	0.59	0.68	0.63
gpt-3.5-turbo (single prompt)	0.66	0.71	0.68
gpt-4 (single prompt)	0.74	0.58	0.65

Table 6: Accuracy of different versions of GPT on the Neuropathic pain dataset.

do not form an edge. We notice that many of the variable names are in Swedish, so we employ an LLM (gpt-4-web) to translate the names to English as a preprocessing step.

Results. Table 6 shows the results of different LLM-based methods. As a baseline, using an algorithm that return either direction at random for each pair, we would obtain

0.5 recall and 0.25 precision on recovering the true edges, leading to an F1 score of 0.33. Using a simple prompt, “A causes B. R and L refer to the right and left sides of the body respectively. Answer with true or false.”, Tu et al. (2023) use chatgpt-3.5-web to obtain an F1 score of 0.21. This LLM performance is worse than the random baseline which may indicate that LLMs are not useful for full graph discovery. However, with a different prompt, we find that the same model can provide over 3X increase in the F1 score. We provide the LLMs with an identical prompt to the “single” prompt from the pairwise task, with one addition: we add a third option, “C: No causal relationship exists”. With this prompt, gpt-3.5-turbo (API analogue of chatgpt-3.5-web) obtains 0.68 F1 score. This score is double the F1 score expected from a random guess baseline, indicating that the output from LLMs provide a non-trivial utility for inferring causal edges.

Algorithm	NHD	No. of predicted edges	Baseline NHD	Ratio
TCDF	0.33	9	0.39	0.84
NOTEARS (Static)	0.33	15	0.44	0.75
NOTEARS (Temporal)	0.35	7	0.38	0.92
DAG-GNN (Static)	0.32	23	0.49	0.65
DAG-GNN (Temporal)	0.34	16	0.44	0.77
gpt-3.5-turbo	0.33	62	0.76	0.43
gpt-4	0.22	46	0.65	0.34

Table 7: Normalized hamming distance (NHD) for different causal discovery algorithms. Since NHD depends on the number of predicted edges, we compare the ratio of NHD and baseline NHD across algorithms. A lower NHD ratio is better. LLM-based discovery (gpt-3.5-turbo) obtains comparable NHD and the lowest NHD ratio compared to recent covariance-based discovery algorithms.

3.2.2 Arctic sea ice dataset

To evaluate LLM’s ability to discover causal graphs, we now consider a dataset from a different field of science: arctic sea ice and atmospheric science. This dataset is on the drivers of arctic sea ice thickness (or coverage): what causes the arctic sea coverage to increase or decrease? In a recent publication, Huang et al. (2021) provide a domain knowledge-based graph with 12 variables and 48 edges. Variables in the graph include total cloud water path, sea level pressure, geopotential height, meridional and zonal wind at 10m, net shortwave and longwave flux at the surface. Importantly, the graph contains some double-sided edges.

Since the number of variables is low (12), we are able to conduct an experiment to discover the full graph using LLMs. Huang et al. (2021) evaluate the accuracy of three recent discovery algorithms: temporal causality discovery framework (Nauta et al., 2019), NOTEARS (Zheng et al., 2018), and DAG-GNN (Yu et al., 2019). They evaluate the algorithms on the normalized Hamming distance (NHD) between the predicted graph G' and the ground-truth graph G . For a graph with m nodes, the distance is given by $\sum_{i,j=1}^m \frac{1}{m^2} 1_{G_{ij} \neq G'_{ij}}$, the number of edges that are present in one graph but not the other, divided by the total number of all possible edges. The distance is zero if G' outputs exactly the same edges as G . For reference, since the number of nodes and edges in our ground-truth graph are 12 and 48 respectively, if G' outputs 48 completely different edges compared to the 48 edges of G , the NHD between G and G' will be 0.66. However, if we output no edges at all, the NHD will be better than the algorithm, 0.33. Since the NHD depends on the number of edges returned by a discovery algorithm, we compare the ratio of NHD of the discovery algorithm and a “floor” baseline that outputs the same number of edges but all of them are incorrect. A lower ratio implies the multiple by which the graph discov-

ery algorithm is better than the worst baseline returning the same number of edges.

Table 7 shows the NHD metric for the three covariance-based causal discovery algorithms and knowledge-based LLMs. For each algorithm, we also show the number of predicted edges and the NHD for a worst case algorithm that predicts the same number of edges. NHD for the three covariance-based algorithms is similar. However, the ratio of NHD for the algorithm and the baseline algorithm provides the relative improvement that the algorithm provides compared to the baseline. Using the ratio metric, we see that DAG-GNN (static) performs the best and has the lowest Ratio distance.

For LLM-based discovery, we use the same prompt as in the neuropathic pain dataset. Using gpt-turbo-3.5, we obtain a normalized hamming distance of 0.33, comparable to the three recently proposed covariance-based causal discovery algorithms. However, the LLM returns 62 edges, so its NHD ratio (0.43) is substantially better than covariance-based algorithms. Finally, gpt-4 obtains the lowest NHD among all algorithms and significantly outperforms state-of-the-art covariance-based discovery algorithms. It outputs almost the same edges as the ground-truth (46) with an NHD of 0.22, one-third less than other algorithms. Its NHD ratio is also the lowest. gpt-4 recovers over half (29) of the edges correctly with a precision of 0.63 (F1 score=0.57), indicating the competitiveness of LLM-based algorithms for causal discovery.

3.3 Probing LLM behavior further

TL;DR: The experiments reveal that GPT-3.5 and GPT-4 have memorized a significant portion of the Tübingen dataset, implying that our benchmark experiment is valid for measuring an LLM’s ability to transform its knowledge into a causal answer, but is not valid for estimating the likelihood an arbitrary relationship has been memorized. Other experiments probe further into semantics and explore what the LLM pays attention to in the prompt.

To better characterize the validity and robustness of our experiments at the intersection of LLMs and causal discovery, we report the results of our LLM probing experiments, described in Sec 2.4.

Memorization Test Results: We provide the LLM with the first 3 columns of the dataset (the row ID and 2 variable names), and ask the model to complete the remaining 2 columns of the row (the source dataset name and the ground truth causal direction. We find that GPT-3.5 is able to correctly recall 58% of the remaining cells in the dataset, and recall 19% of the rows without error. Our memorization test with GPT-4 is able to recall 61% of remaining cells in the dataset and 25% of rows without error. These results indicate that the Tübingen dataset is certainly in GPT’s training set and that GPT has had an opportunity to memorize the dataset. However, the gap between percentage of inputs memorized and LLMs accuracy is still significant and more work is needed to understand the role of memorization for the inputs that are not captured by our memorization test.³

More generally, we expect knowledge-based causal discovery to be founded partly on memorized information, but also dependent on the LLM’s ability to process and transform seen causal relationships for use in multiple contexts. We can conceptually model the LLM’s end-to-end efficacy $P(Y|D)P(D)$, where $P(D)$ represents the likelihood of a causal relationship being memorized by the LLM, and $P(Y|D)$ the likelihood that the relationships can be correctly transformed. Knowing that our benchmark dataset is memorized implies that our experiments can only measuring $P(Y|D)$, but not measure $P(D)$. That is, when evaluating an LLM on an existing benchmark, its performance on the benchmark is a test of how well the LLM’s ability to process and transform its knowledge into the necessary causal relationship assuming that it has been trained using some representation of the underlying information.

³The neuropathic pain and the arctic sea ice coverage dataset were not made available online in a text file format and hence are less amenable to a straightforward memorization test.

Which cause-and-effect relationship is more likely? A. changing SLOT1 causes a change in SLOT2. B. changing SLOT3 causes a change in SLOT4. Let’s work this out in a step by step way to be sure that we have the right answer. Then provide your final answer within the tags <Answer>A/B</Answer>.

Figure 4: We probe the importance of individual words for getting a correct answer by redacting a random word from a question. Here, we show our results for experiments in gpt-3.5-turbo, averaged across 357 random redaction probes over our Tübingen experiment. We highlight words based on their importance for getting a correct result. A white background indicates that redacting the word did not reduce accuracy. A dark blue highlight indicates that redacting the word reduced accuracy the most.

Redaction Test: To characterize whether or not the LLM is attending to the correct words and concepts in a question when answering our pair-wise causal discovery questions, we use a series of redaction and perturbation tests. Our redaction test shows what words are important for correctly labeling a causal relationship between two features. See Figure 4. If the LLM is correctly reading the question, we would expect that it will pay a lot of attention to key instructions and to the options themselves.

4 LLMs for Actual Causality and Causal Judgments

TL;DR: The ability of LLMs like GPT-4 to capture and apply common sense and domain knowledge enables substantial improvements in counterfactual reasoning and actual causality tasks, making them valuable in real-world applications. However, they can still struggle with ambiguity and exhibit unpredictable failure modes, requiring additional context for accurate answers.

In the previous section, we saw how LLM’s ability to capture common sense and domain knowledge allows for knowledge-based approach to causal discovery. In this section, we now study how LLMs’ can bring these capabilities to enable a systematized approach to actual causality challenges. Actual causality is motivated by problems of attribution and assigning responsibility in real world applications, such as legal reasoning, machine failure debugging, and root-cause analysis for system regressions.

For example in tort law, the core problem of deciding how much blame for a plaintiff’s injury is attributed to a defendant’s action relative to other potentially mitigating events is fundamentally an actual causal inference task.

Causal inference researchers have attempted to use structural causal models (SCMs) and other formal causal models to define actual causality in a way that is consistent with how humans naturally attribute cause and related concepts of responsibility, blame, and explanations to events and their outcomes (Halpern, 2016; Kueffner, 2021). This task has proven exceedingly difficult, because human judgments of causality depend on elements of background context that are difficult to formalize in an SCM. Types of background context that are fundamental to human causal judgments but difficult to formalize in an SCM include:

- **Causal frame:** The set of candidate causal events that are relevant to a particular outcome event. SCMs require relevant causes be included as endogenous variables, but in practice humans rely on domain knowledge and common sense to set the causal frame after the outcome occurs (Icard & Goodman, 2015). For example, there are a multitude of potential causes of forest fires, but in the case of a particular forest fire, upon learning lightning struck and that a hiker was carelessly smoking a cigarette during fire season, we know these are relevant causal events. We would also know whether to ignore “drought” or “lack of controlled burn” if they are not relevant in the case of this particular fire.
- **Necessary causality:** Whether a candidate cause needed to happen for the outcome to occur.
- **Sufficient causality:** Whether a candidate cause’s occurrence would have led to the outcome event if other causal events had occurred differently.
- **Normality:** The degree to which causal events align with statistical norms or prescriptive norms (social, moral, or legal norms) (Phillips et al., 2015; Kominsky et al., 2015). When agents violate norms, they are typically judged to be more of a cause of resulting outcomes (Phillips et al., 2015; Kominsky et al., 2015). Human causal judgments depend highly on whether candidate cause is a norm violation, or whether it is more or less normal than other causal events.
- **Other human factors:** Other human factors include bias towards action, handling intention and epistemic state, and how bad outcomes are interpreted. When the candidate cause is an agent’s behavior, humans tend to ascribe more causality actions (e.g., tossing a

cigarette) than to lack of actions (e.g, not doing a controlled burn) (Henne et al., 2017). In addition, when the candidate cause is an agent’s behavior, whether the agent acted with intention and knew what they were doing (Nadelhoffer, 2006; Knobe, 2003) matters. Finally, human causal judgments also depend on whether the outcome is undesirable (e.g., causing a forest fire vs. causing a reforestation initiative) (Kominsky et al., 2015).

Sidestepping the challenge of formalizing these concepts into a causal model, LLMs offer an opportunity to capture the necessary and relevant background context for an event directly from natural language description of an event. Given that an LLM is trained on text narratives written by humans, the subjective elements of causal judgments may be encoded as part of the LLM’s internal representations.

In this section, we first review the concept of counterfactual reasoning and show that LLMs can answer counterfactual questions about an event with accuracy close to that of humans on a benchmark task. We then present an example scenario to illustrate the challenges of modeling real-world events using actual causality and introduce the two key counterfactual concepts used in formal models of actual causality: necessity and sufficiency. Using 15 different vignettes commonly used in actual causality literature, we provide evidence that LLM can capture the necessary background context and reason about necessity and sufficiency of actual causes. Finally, we show how this capability transfers to a causal judgment benchmark task from Big Bench (Suzgun et al., 2022), specifically designed to test LLMs on their ability to infer the normality of a candidate causes.

4.1 Building block of actual causality: Counterfactual reasoning

TL;DR: LLMs like GPT-3.5 and GPT-4 have shown substantial improvements in counterfactual reasoning ability, with GPT-4 achieving 92.44% accuracy, just six percentage points below the human baseline.

Counterfactual reasoning is the process of considering hypothetical situations or alternate realities by altering specific elements or conditions of an actual event or situation (Kahneman & Miller, 1986; Byrne, 2005). Such reasoning is a key element of actual causality. To determine the causes of an event, it is important to simulate alternative worlds where an event may not have happened and reason about the consequences. For example, a naive approach to

Premise	Counterfactual Question	Multiple-choices answers
A woman does not order Chinese food.	What would have happened if she had ordered Chinese food?	The woman would have become less hungry.;The woman would have become very hungry.;That is not possible.
A woman sees a fire.	What would have happened if the woman had touched the fire?	She would have been burned.;She would not have been burned.;That is not possible.;She would have seen fire.
A bird lands in a forest.	What would have happened if a plane had landed in the forest?	The plane would have crashed.;Everything would have been fine.;The plane would have landed safe and sound.;In a forest you will find lots of planes.
A plant grows in a planter.	What would have happened if the planter grows in the plant?	That is not possible.;It would have grown quicker.;The plant would have suffered.;The planter would have cultivated the plant.
A mortician prepares a corpse.	What would have happened if the mortician had prepared a dinner?	He would have had a delicious dish.;Morticians cannot prepare dinners.;The dinner would have been buried.;The mortician would have killed the corpse.
An oil tanker sails across an ocean.	What would have happened if the oil tanker had broken up in an ocean?	There would have been environmental pollution. That is not possible.;The oil tanker would have continued to carry oil.;The oil tanker would have been saved.;
A car crashes into a tree.	What would have happened if the car had parked beneath the tree?	Nothing special would have happened.;The car would have been hit by the tree.;That is not possible.;I think it would have crashed into the tree.
A child draws a picture.	What would have happened if the child had erased the picture?	The picture would not have been visible.;The picture would have been visible.;That is not possible.
A craftsman builds a house.	What would have happened if the house had built a craftsman?	That is not possible.;The house would have been built faster.;Everything would have been fine.;The craftsman would have hands.
A doctor washes their hands at work.	What would have happened if the doctor hadn't washed their hands?	The patients could get an infection.;The patients could get better.;That is not possible.

Table 8: Example scenarios from the CRASS counterfactual reasoning benchmark. The task is to select the best answer choice for the counterfactual question, given a premise.

Model	Accuracy
GPT-3 (Froberg & Binder, 2022)	58.39
T0pp (Sanh et al., 2021)	72.63
text-davinci-003	83.94
gpt-3.5-turbo	87.95
gpt-4	92.44
Human annotators	98.18

Table 9: Accuracy of different LLMs on the CRASS counterfactual reasoning benchmark. gpt-4 achieves 92% accuracy, significantly higher than the previous reported accuracy on this benchmark and within six percentage points of human annotators’ accuracy.

actual causality is to use a definition based on *counterfactual dependence*: an event A is a cause of another event B if B would not have happened without A.

Independent of actual causality, counterfactual reasoning is a desirable skill for a LLM, as it would assist users in decision-making, planning, and offer insights that may not explicitly apparent from the original context. With this motivation, Froberg & Binder (2022) provide a benchmark called CRASS (Counterfactual Reasoning Assessment) for evaluating the ability of LLMs to answer counterfactual

conditional questions, a subset of which is included in the BIGBench collection of datasets (Srivastava et al., 2022).⁴ The benchmark contains 275 instances, where each instance has the LLM select from multiple choice answers to a counterfactual conditional question such as “A woman opens a treasure chest. What would have happened if the woman had not opened the treasure chest?”. Table 8 lists some instances from the dataset.

The authors report GPT test-davinci-003 as the best performing (zero-shot) LLM on this task, with a top-1 accuracy of 58.39%. Another language model (T0pp) (Sanh et al., 2021), which was finetuned on a multi-task setup of classification tasks, obtains 72% accuracy. This is contrasted with 98.18% accuracy human baseline, indicating that counterfactual predictions remain a challenging task for LLMs.

We revisit the claim using LLMs released after the benchmark’s publication. As shown in Table 8, each problem instance has a premise, counterfactual question and a list of answer options. We construct a prompt for the LLM by concatenating the premise and the counterfactual

⁴The full dataset was released on Github on August 2021, so it is after the training cutoff for text-davinci-003 model. But it is a month before the cutoff for gpt-3.5-turbo. A smaller subset of inputs was released as a part of BIG-Bench in May 2021 with a canary string for LLMs to avoid including it in their training data.

question, and then presenting the answer options as A, B, C, D (for an example prompt, see Suppl. B.1). For gpt-3.5-turbo and gpt-4, we additionally provided a system message, “You are a helpful assistant for counterfactual reasoning.”

Table 9 shows the accuracy of different LLMs. Remarkably, GPT 3.5 version models show substantial improvement in accuracy over GPT 3. gpt-3.5-turbo obtains an accuracy of 87.95. gpt-4 improves it further to 92.44%, which is 20 percentage points higher than the previous best accuracy. Comparing to the human accuracy on this task, gpt-4 is within six percentage points of average human accuracy. The results indicate that large language models of GPT-3.5 and 4 series represent a substantial jump in the counterfactual reasoning ability.

Understanding LLM output. We now try to analyze some outputs from gpt-4 to understand its reasoning capabilities. For full LLM outputs of the prompt discussed here, see Suppl. B.1. As the high accuracy indicates, gpt-4 is able to simulate different scenarios based on the text prompt and answer what-if scenarios. In many cases, when gpt-4’s answer does not match the benchmark, it reveals ambiguity in the text used to describe the scenario. In one of the cases, the LLM is asked what will happen if a man catches a water balloon, given that he has caught a ball. The benchmark answer is that he will get wet whereas the LLM output correctly argues that he may get wet or not, depending on whether the balloon bursts as he catches it. Similarly, another question asks what will happen if a man gets very nervous. The benchmark answer indicates that he will pass out, but the LLM output correctly mentions that, “*Passing out due to extreme nervousness is possible, but not guaranteed.*” and concludes that the man would not pass out. We suspect that given enough clarifying context, the LLM should be able to correctly answer such questions.

That said, there are a few instances where the LLM does not capture context that people may easily determine through common sense. For instance, one of the questions asks what would happen if a man walks on a bed instead of walking on the street. “Walking on a bed” conveys that the man is likely indoors, but the LLM incorrectly imagines a scenario where the man is still walking to the same destination, but now he is walking over a bed, which makes it incorrectly conclude that the man “will get late” (see Suppl. B.1). These results indicate the importance of minimizing ambiguity when interacting with an LLM—wherever possible, adding small additional details (e.g., “walking on his bed at home”) can increase the chances of a correct answer.

4.2 Inferring actual causality: Necessary and sufficient causes

TL;DR: LLMs, specifically GPT-4, show significant capability in understanding scenarios in natural language and identifying necessary and sufficient causes based on counterfactual reasoning. This capability of reasoning directly over text was not possible before LLMs.

Given the promising performance of LLMs on a counterfactual reasoning task, we now investigate whether LLMs can infer whether events are necessary or sufficient for a focal event.

4.2.1 Motivating Example: Necessary and sufficient causes

To motivate, we present the following example story of a bottle of beer spilling at a party, a variation on the story presented in Henne & O’Neill (2022). Since the story was published in 2022, it cannot have been in the training set of either GPT 3 or GPT 4 models.

PROMPT: Mike, Jack, Susan, and Peter are at a party. There is an open bottle of beer on the table. Mike is moving through the room and accidentally bumps against the table. Jack saw that the bottle was about to fall, and he reached out to catch the bottle. But just as Jack was about to catch the bottle, Peter bumped into Jack, which caused Jack to miss the bottle, which then fell to the ground. As a result, the bottle of beer fell to the floor, spilling the beer. Right after the bottle hit the floor, Susan tripped and fell onto the table.

Let us unwrap the story progressively to bring out the two key concepts of actual causality: *necessity* and *sufficiency*. **Necessary causality:** The property of an event, C, that if it had not occurred, then the outcome event, E, would not have occurred.

So we start by evaluating whether the LLM can make simple statements about necessary causality in the context of the beer bottle spilling story. First start by ask the LLM to provide a one sentence counterfactual explanation as to why Mike caused the spill.

PROMPT: Mike, Jack, Ben, Susan, and Peter are at a party. There was an open bottle of beer on the table. Mike is moving through the room and accidentally bumps against the table. As a result, the bottle of beer fell to the floor, spilling the beer. Did Mike cause the beer to spill? Provide a one sentence counterfactual explanation as to why or why not.

Limiting the response to one sentence discourages more verbose responses.

LLM: Yes, Mike caused the beer to spill; if Mike hadn't bumped into the table, the beer bottle would not have fallen and spilled.

The LLM zeroes in on the necessary causality as the explanation. More generally, we can prompt the LLM to invoke the necessary cause principle and output the actual cause for an event.

In the above example, Mike's bumping into the table was necessary for the beer to spill. Moreover, it was also *sufficient* to cause the beer to spill. But there are many cases where a causal event is a necessary but insufficient cause of an outcome.

Sufficient causality: The property of an event, C, that if it had occurred, then the outcome event, E, would have occurred.

Formal models provide unambiguous definitions of sufficient causality; in Pearl's counterfactual notation, the probability a causal event $C=1$ is a sufficient cause is $P(E_{C=1} = 1 | C = 0, E = 0)$ (Pearl, 2009b). In natural language this reads as "in cases where C was 0 and E was 0, what is the probability that E would have been 1 if C had been 1?"

But this phrasing awkward to express in situations where C and E were both indeed 1, not 0. So we turn to cognitive science literature to find a more natural text expression of sufficient cause. We use a variant of *robust sufficient causality* (Hitchcock, 2012; Woodward, 2006; Hitchcock, 2012):

Robust sufficient causality: Given that causal event C occurred, outcome event E would have occurred even if other events that contributed to B's occurrence had not occurred. To elicit a statement about robust sufficiency, we prompt the LLM to state what other causal events lead to the spill, and then predict a counterfactual about whether the spill would have occurred those events had not happened. This is a more challenging task than the necessary cause prediction, because it requires the LLM to solve a *causal frame problem* decide which other causal events to, some of which may their interactions.

PROMPT: Based on the information in the passage, what other events directly contributed to the fall of the beer and what would have happened if those events hadn't happened?

LLM: Other than Mike bumping into the table, the presence of the open beer bottle on the table directly contributed to the fall of the beer; if the bottle had not been placed on the table or had been properly closed, it would not have fallen and spilled even if Mike had bumped into the table.

This response highlights the capability of the LLM to set the *causal frame*, i.e., select causal events relevant to the

inference problem. We did not specify which causal events were eligible for consideration. For example, the LLM could have included the presence of the table as a causal event, but it did not. This is in contrast to formal models such as SCMs where all the variables must be articulated in advance of running an inference query on the model. We conjecture that, as a probability model, the LLM is selecting probable causal abstractions based on both the prompt and statistical regularities in the training data. This allows its reasoning to adapt based on the scenario; if Mike bumping the table is the only causal event in consideration, it is sufficient, but that may change if other causal events are present. This highlights the flexibility of LLM's in causal analysis relative to SCMs and other formal models.

4.2.2 Evaluating LLMs on inferring necessary and sufficient causes

When the principle of necessity and sufficiency yields the same cause, actual cause of an event can be easily determined. However, in many real-world scenarios, they may output different answers. For example, in the beer spilling example above, suppose there were two beer bottles on the table. Then the presence of any beer bottle on the table is a sufficient cause for beer spilling, but not necessary. Even if one of the bottles was not placed on the table, the other one would have spilled. Therefore, inferring necessary and sufficient causes typically requires applying the definitions of necessity and sufficiency and using formal reasoning (Halpern, 2016).

Below we investigate whether LLMs can provide this reasoning directly using the natural language description of an event. We evaluate on a collection of 15 vignettes (example scenarios) from Kueffner (2021). Table 10 shows sample vignettes. These vignettes are widely used to discuss and critique actual causality definitions in the literature and span challenging scenarios across seven different types: symmetric overdetermination, switch, late preemption, early preemption, double preemption, bogus preemption, short circuit, and other miscellaneous examples. For example, symmetric overdetermination is a scenario where multiple processes (causes), all of which producing the same outcome (event), terminate at the same time. Suppl. C.1 provides a description for each type.

For each vignette, our goal is check whether an LLM can correctly identify whether a cause is necessary, sufficient, or both. To do so, we use the following prompt template, where PRINCIPLE is replaced by either "minimal change" (corresponding to necessary cause) or "multiple sufficient causes" (corresponding to sufficient causes). In each example, we provide two options: one of them is the necessary cause and the other the sufficient cause. Interestingly, the

Vignette Type	Input Context	Event	Actor	Nec.	Suff.
Overdetermination	Alice (AF) and Bob (BF) each fire a bullet at a window, simultaneously striking the window, shattering it (WS).	window shattering	Alice	No	Yes
Switch	Alice pushes Bob. Therefore, Bob is hit by a truck. Bob dies. Otherwise, Bob would have been hit by a bus, which would have killed him as well.	Bob's death	Alice	No	Yes
Late preemption	Alice (AF) and Bob (BF) each fire a bullet at a window. Alice's bullet hits the window first (AH). The window shatters (WS). Bob's bullet arrives second and does not hit the window (BH).	window shattering	Alice	No	Yes
Early preemption	Suppose Alice reaches out and catches a passing cricket ball. The next thing on the ball's trajectory was a solid brick wall that would have stopped the ball. Beyond that there was a window.	window being intact	Alice	No	Yes
Double preemption	Alice intends to fire a bullet at a window (AI). Bob intends to prevent Alice from hitting the window (BI). Bob tries to stop Alice (BSA). Bob is stopped by Carol (CSB). Alice fires a bullet (AF), hits the window (AH) and shatters it (WS). The window shatters (WS).	window shattering	Alice	Yes	No
Bogus preemption	Alice intends to put lethal poison into Carol's water. However, Alice does not put lethal poison into Carol's water (\neg AP). Bob puts an antidote into Carol's water (BA). The water is lethal (L), if the poison is added without the addition of an antidote. If Carol would consume the lethal water she would die (CD). Carol consumes her water (CC). Carol does not die (\neg CD).	Carol's survival	Alice	No	Yes
Short circuit	Carol is alive (CA). Alice puts a harmless antidote in Carol's water (AA). Adding an idote to the water, protects it against poison (WS - 'water save'). If Alice puts the antidote into Carol's water, Bob will poison the water (BP) Adding poison to an unprotected water makes it toxic (WT). If Carol would drink toxic water she would die (i.e. inhibiting CS). Carol consumes her water and survives (CS).	Carol's survival	Alice	No	Yes
Miscellaneous	If there is hot weather, flowers will die. Watering prevents the flowers to die in hot weather. The neighbor does not water the flowers in her yard. The flowers die.	flowers' death	neighbor's inaction	Yes	Yes

Table 10: Example vignettes for evaluation of inferring necessary and sufficient causes, categorized by their *type* based on the different ways in which potential causes can interact to yield the final outcome. Each vignette tests two questions: “*Is {Actor} a necessary cause of {Event}?*” and “*Is {Actor} a sufficient cause of {Event}?*”.

prompt itself was constructed using the gpt-3.5-turbo (see Suppl. C.2 for details). An example prompt for the early preemption vignette is shown below.

SYSTEM: You are an expert in counterfactual reasoning.
Given an event, use the principle of [PRINCIPLE] to answer the following question.
PROMPT: Suppose Alice reaches out and catches a passing cricket ball. The next thing on the ball's trajectory was a solid brick wall that would have stopped the ball. Beyond that there was a window.
Is Alice a [NECESSARY/SUFFICIENT] cause of window being intact?
After your reasoning, provide the final answer within the tags <Answer>Yes/No</Answer>.

Using this prompt, we evaluate all 15 example scenarios on two LLMs: gpt-3.5-turbo and GPT-4. For example, on this vignette, both LLMs correctly answer that Alice is not a necessary cause; even without her action, the brick wall would have stopped the ball. For evaluating sufficiency, gpt-4 correctly answers that Alice is a sufficient cause of the window being intact, but gpt-3.5-turbo answers incorrectly (see responses in Suppl. C.4.1).

Table 12 summarizes the results over all vignettes. We see a significant difference between the accuracy of gpt3.5-

turbo and gpt-4, indicating a marked change in ability between the two models. gpt-3.5-turbo fails to capture the nuances of necessity and sufficiency definitions and obtains an accuracy near random guess. However, gpt-4 is accurate for most vignette types. This is remarkable because the LLM was not explicitly provided definitions of necessary and sufficient causes. As with causal discovery, however, the high accuracy also comes with unpredictable failure modes. For example, even as gpt-4 use the correct sufficiency principle for most vignettes, on the short circuit vignette (see Table 10), gpt-4 uses only counterfactual dependence principle to answer the question and fails to reason about the sufficiency of cause. The input prompt and incorrect response are shown in Suppl. C.4.2.

Since the above vignettes or their variants have been quoted in many publications, there is a concern that LLMs may have memorized them and may not be reasoning about the answers. We therefore report results on the “lab-vignettes” dataset, a new dataset we created that adapts the base vignettes to a laboratory setting. We do so by keeping the same type of scenario (e.g., short circuit), but changing the situation to that in a laboratory with reagents, test tubes, mixtures and crystals. Suppl. C.3 provides details on this dataset. On this new, unseen dataset, we find a similar pat-

Vignette Type	Necessary	Sufficient
<i>gpt-3.5-turbo</i>		
Overdetermination	✓, ✓	X, ✓
Switch	X,X	✓, X
Late preemption	X	X
Early preemption	X, ✓, X	X, X, ✓
Double preemption	✓	✓
Bogus preemption	✓	X,
Short circuit	X	X
Miscellaneous	X, ✓, ✓, X	✓, ✓, X, ✓
Total Accuracy	46.6%	46.6%
<i>gpt-4</i>		
Overdetermination	✓, ✓	✓, ✓
Switch	✓, ✓	✓, ✓
Late preemption	✓	✓
Early preemption	✓, ✓, ✓	✓,
Double preemption	✓	X
Bogus preemption	✓	✓
Short circuit	X	X
Miscellaneous	✓, X, ✓, ✓	✓, ✓, ✓, ✓
Total Accuracy	86.6%	86.6%

Table 11: Accuracy of gpt-3.5-turbo and gpt-4 on inferring necessary or sufficient cause on 15 standard vignettes. The vignettes are divided into eight types (e.g., Early Preemption type has three vignettes). Each (✓/X) corresponds to a correct/incorrect answer on a single vignette. gpt-3.5-turbo fails at the task (worse than random chance) but gpt-4 can infer necessary and sufficient cause with high accuracy.

tern when comparing models. gpt-4 obtains significantly higher accuracy than gpt-3.5-turbo. While gpt-4 retains its overall accuracy, we find an interesting pattern: for both models, necessity is an easier concept to answer than sufficiency. gpt-4 obtains over 92% accuracy on deciding necessity of cause compared to 78% on deciding its sufficiency. This may be because necessity always involves comparing the output under a counterfactual world where only the cause is flipped; whereas sufficiency can be nuanced, involving counterfactual flips to all relevant other variables (and deciding what those variables should be).

Overall, our results indicate the capability of LLMs to understand a scenario in natural language and output actual causes according to a pre-specified definition, but also show the lack of robustness due to unpredictable failures.

Vignette Type	Necessary	Sufficient
<i>gpt-3.5-turbo</i>		
Overdetermination	✓, ✓	X, ✓
Switch	X, ✓	✓, X
Late preemption	X	✓
Early preemption	✓, X	X, X
Double preemption	✓	✓
Bogus preemption	✓	X
Short circuit	X	X
Miscellaneous	✓, ✓, ✓, X	✓, X, X, ✓
Total Accuracy	64.2%	42.8%
<i>gpt-4</i>		
Overdetermination	✓, ✓	✓, ✓
Switch	✓, ✓	X, ✓
Late preemption	✓	✓
Early preemption	✓, ✓	X
Double preemption	✓	✓
Bogus preemption	✓	✓
Short circuit	✓	✓
Miscellaneous	✓, X, ✓, ✓	✓, ✓, X, ✓
Total Accuracy	92.8%	78.5%

Table 12: Testing dataset memorization issues with a novel “lab-vignettes” dataset. The average accuracy of gpt-4 stays the same as in the std vignettes, indicating that gpt-4’s capabilities to infer necessary and sufficient cause can generalize to new data. Inferring necessary cause (93%) emerges as an easier task than inferring sufficient cause (78%).

4.2.3 Assessing responsibility

LLM-based causal reasoning can allow for a more nuanced understanding of causality, taking into account factors such as intention, moral obligation, and epistemic state that are difficult to capture using formal SCM approaches. We present a study of these capabilities in Supplementary D.

4.3 Evaluating LLMs ability to infer normality

TL;DR: While large language models have performed less well on the causal judgment benchmark, they show promise in answering the core common sense elements needed to answer these questions. One of these elements is *normality* - e.g., whether social norms are violated.

Next, we look at examples of stories from the causal judgment benchmark in BIG-Bench Hard (Suzgun et al.,

2022). The causal judgment benchmark collects examples from vignettes similar to those in Kueffner (2021), but are in more general language and provide more background information. Each of the stories ends with a yes/no question about whether a particular event caused the outcome. The benchmark provides ground truth labels for each question, collected from human annotators. It is an ideal benchmark for evaluating how well LLMs would do in addressing practical actual causality problems, as opposed to more contrived examples.

The benchmark has proven resilient to saturation as new generations of LLMs are released. For example, while several benchmarks in BIG-Bench Hard saw a dramatic increase in accuracy between GPT3 and GPT4, causal judgment had a small increase (61.5% in text-davinci-003 and has an accuracy of 67.6% in GPT-4-32K (AML-BABEL, accessed 2023-04-25). Top human performance is at 100% accuracy.

In a preliminary analysis, for each passage in the benchmark, we manually labeled whether the causal event in question was “normal” or “abnormal.” We then evaluate how well LLMs can classify normality.

Constructing a benchmark for normality. First, we filter examples that ask whether an agent intentionally caused an outcome and only focus on examples that ask if an event or agent caused an outcome. We do this because assessing intention is a separate task from ascribing causality, and because many intention examples were contrived (e.g., trolley problems) such that normality does not apply. We then manually label each example as normal or abnormal. In most cases these labels were not subjective assessments on our part. Each element in the causal judgment data set includes a citation to the source of the passage. In many of these studies, normality was explicitly varied across sections of a vignette, and thus we could confirm our labels from their source papers. However, in some cases, we had to make subjective best-guess assessments.

Prompting the LLM. We follow a two-step process. We first prompt the LLM to extract the causal event in question.

PROMPT: The last sentence of the following passages is a question. The question either ask if an outcome event happened because a certain causal event occurred. When the question asks specifically about “intention,” the causal event is not the outcome, it is the action that lead to the outcome. State the causal event being asked about. Do not include the outcome event in the statement. Do not include other causal events in the passage.

Passage: The receptionist in the philosophy department keeps her desk stocked with pens. The administrative assistants are allowed to take the pens, but faculty members are supposed to buy their own. The administrative assistants typically do take the pens. Unfortunately, so do the faculty members. The receptionist has repeatedly emailed them reminders that only administrative assistants are allowed to take the pens. On Monday morning, one of the administrative assistants encounters Professor Smith walking past the receptionist’s desk. Both take pens. Later that day, the receptionist needs to take an important message... but she has a problem. There are no pens left on her desk. Did Professor Smith cause the problem?

LLM: Professor Smith took pens from the receptionist’s desk.

We then ask the LLM to evaluate the normality of the causal event in question. Inspired by both model-based and subjective notions of normality, we construct the following prompt.

PROMPT: Given a causal event that causes an outcome event, the causal event is “abnormal” if any of the following is true:

- The causal event’s occurrence was unexpected, unlikely, surprising, rare, or improbable
- The causal event was an agent’s action or lack of action that non-accidentally, knowingly, or negligently violated social, legal, or ethical norms.

The opposite of “abnormal” is “normal”. A causal event is “normal” if its occurrence was expected, likely, probable, and not surprising. In cases where the causal event is an agent’s action or lack of action, the action or lack of action can also be considered “normal” if it intentionally compliant with social, legal and ethical norms. The normality of a causal event is independent of whether the outcome event was bad or and whether the outcome event was normal.

Please read the following passage and statement about a causal event from the passage. Provide your step-by-step rationale as to why the causal event is abnormal or normal. Then, return the word “normal” if you think the causal event is normal or “abnormal” if you think the causal event is abnormal.

Passage: [PASSAGE]

Event: Professor Smith took pens from the receptionist’s desk.

LLM: Professor Smith’s action of taking a pen from the receptionist’s desk is abnormal because it violates the established norm that faculty members are not allowed to take pens from the receptionist’s desk. Despite the receptionist sending reminders about this rule, Professor Smith still took a pen, making his action unexpected and improper.

In defining normality, we attempted to capture both statistical norms (with keywords “expected,” “likely,” “probable,” “frequent,” “unsurprising,” and their opposites) and prescriptive norms. Violation of prescriptive norms can depend on intention, so we use keywords “nonaccidentally,” “knowingly,” “negligently,” and “intentionally.” We also ask the model to provide a chain-of-thought “step-by-step” rationale for its answer.

Results. Initial results show accuracy in inferring normality was 66.5% on text-davinci-003 and 70.2% on GPT-4-32K, which are comparable to performance on the core benchmark task. Table 13 shows three examples of correct and incorrect answers across answers by text-davinci-003 and GPT-4-32K.

Table 13 demonstrates the capability of LLMs to parse normality but also provides intuition for how the LLMs can get normality wrong. While results did improve between versions of GPT3 and GPT4, there were some cases where GPT4 was wrong when GPT3 was right, although its rationale when it was wrong was typically more persuasive.

4.4 Discussion

SCMs and other formal models of actual causality have struggled to represent background common sense elements of a causal judgment, such as necessity, sufficiency, normality and responsibility. We show through analysis of the CRASS (Frohberg & Binder, 2022) and causal judgment (Ghazal et al., 2017) benchmarks and related vignettes from Kueffner (2021) that LLMs do well in reasoning with these elements. Relatively lower performance for direct and chain-of-thought answering on the causal judgment benchmark highlight an opportunity to research and engineer systems that guide LLMs in assembling these core elements into solutions for practical problems in actual causality questions.

5 A New Frontier for Causality

TL;DR: LLMs can augment human expertise in causal analyses with implications for both practice and research. They have the potential to enable end-to-end pipelines for causal inference and facilitate fluid conversational interfaces. Important research opportunities include continuing to explore their capabilities in various causal tasks, human-LLM collaboration, and understanding and improving causal reasoning.

Our results on evaluating large language models on causal discovery, counterfactual reasoning, and actual causality demonstrates that they bring significant new capabilities across a wide range of causal tasks. We believe that these capabilities signal the beginnings of a new frontier and will transform the practice and research of causality. We start by summarizing key takeaways on how to interpret the new capabilities that LLM bring to causality. We then provide implications for practitioners and future research directions at the intersection of LLMs and causality.

What is new with causality and LLMs?

- LLMs provide access to domain knowledge that was, heretofore, only available via human domain experts. LLMs can provide this domain knowledge when explicitly asked, and also implicitly, e.g., as we see in background knowledge necessary for actual causality vignettes.
- LLMs provide a flexible, natural language-based interaction mechanism for conducting causal analysis that can work alongside existing tools, thus democratizing access to causal tools and knowledge.
- LLMs offer a new capability to extract the key primitives of an actual causality question (necessity, sufficiency, normality, etc.). This opens up the possibility of building systems based on formal actual causality research for practical scenarios for perhaps the first time.

It is important to recognize that answering causal questions require iterative steps in practice. LLMs can potentially enable a wide and effective adoption of causality in practice and break the barriers between different research communities related to causality.

What is not changing with causality and LLMs?

- Especially for high-risk and high-value tasks, people have long relied on rigorous analyses in various formal frameworks to ensure correctness of their decision-making. The need for rigorous, well-documented, and verifiable analyses has not changed. This implies that LLM-driven applications must have tight integration with more formal approaches to causal reasoning. LLMs should be able to fluidly connect their language generation with the use of formal tools, including re-framing or translating natural language scenarios to formal descriptions, and ensuring their natural language responses are in accordance with the conclusions of formal analyses.

Passage and Question	Correct answer	An-	text-davinci-003 rationale	GPT-4-32K rationale
Lauren and Jane work for the same company. They each need to use a computer for work sometimes. Unfortunately, the computer isn't very powerful. If two people are logged on at the same time, it usually crashes. So the company decided to institute an official policy. It declared that Lauren would be the only one permitted to use the computer in the mornings and that Jane would be the only one permitted to use the computer in the afternoons. As expected, Lauren logged on the computer the next day at 9:00 am. But Jane decided to disobey the official policy. She also logged on at 9:00 am. The computer crashed immediately. Did Jane cause the computer to crash?	Jane logging in at 9:00 am is an abnormal event because she followed policy.		Logging on the computer at 9:00 am is consistent with the official policy, making this causal event normal. (incorrect)	Jane's decision to log onto the computer at 9:00 am despite the official policy is abnormal, as it intentionally violates the company's policy put in place to prevent computer crashes. This action is unexpected and goes against the established social and workplace rules. (correct)
The receptionist in the philosophy department keeps her desk stocked with pens. The administrative assistants are allowed to take the pens, but faculty members are supposed to buy their own. The administrative assistants typically do take the pens. Unfortunately, so do the faculty members. The receptionist has repeatedly emailed them reminders that only administrative assistants are allowed to take the pens. On Monday morning, one of the administrative assistants encounters Professor Smith walking past the receptionist's desk. Both take pens. Later that day, the receptionist needs to take an important message... but she has a problem. There are no pens left on her desk. Did the administrative assistant cause the problem?	The administrative assistant taking the pen is a normal event because she followed policy.		Taking pens is expected behavior among both administrative assistants and faculty members according to the rules set by the receptionist. Therefore, the causal event of both the administrative assistant and Professor Smith taking pens is normal. (correct)	Taking a pen from the receptionist's desk is abnormal for the administrative assistant in this situation because they contributed to the problem of the receptionist having no pens left, even though it is typically allowed for them to take pens. However, it is important to note that the faculty member's action of taking a pen is even more abnormal and against the established rules. (incorrect)
Joe was about to go pick up his child from a school in the suburbs. In the parking lot, he stopped to talk with friends. Because of this delay in picking up his son, Joe's neighbor, who waited 15 min for Joe to show up, took the child home in his car. On the way home, the neighbor's car was struck by a drunk driver. Joe's son, the only seriously hurt victim, received severe leg injuries. Did Joe cause his son's injury?	Joe delaying in picking up his son to chat with friends is abnormal because it is a social norm violation.		Stopping to talk with friends in a parking lot is a common and expected behavior, making this causal event quite normal. (incorrect)	Joe stopping to talk with friends in the parking lot is a normal event as people often engage in casual conversations when they encounter acquaintances. It is not unexpected, unlikely, or a violation of social norms. (incorrect)

Table 13: Comparative assessments of normality between text-davinci-003 and GPT-4-32K. Stories taken from the BIG-Bench causal judgments task.

5.1 Implications for Causality Practitioners

TL;DR: LLMs can augment human expertise in causal analyses by assisting in graph creation, validation, and robustness checks, while ensuring trustworthiness and human verifiability. In combination with causal tools, LLMs can enable an end-to-end pipeline for causal inference and facilitate fluid conversational interfaces that merge covariance- and logic-based reasoning.

Augmenting human expertise with LLMs. LLMs can enable effective human-AI interaction that reduces human burden during causal analyses, while ensuring trustworthiness and human verifiability. One of the key reasons causal methods are not widely deployed is because using them requires expertise in expressing the causal assumptions formally through a graph and then verifying them through robustness checks. LLMs can act as an assistant in both these processes, augmenting the capability of a practitioner and helping reduce friction in starting a causal analysis.

Given the challenging task of automated causal discovery, most practitioners start their causal analysis by constructing a graph manually based on their domain knowledge. Rather than creating a graph from scratch, we recommend that practitioners input their dataset metadata to an LLM and obtain an initial version of the graph that they can iterate on. At least for medium sized graphs, our results indicate that LLMs can match state-of-the-art accuracy in learning real-world graphs, as long as there is descriptive metadata available for each variable. Doing so can save time since editing a graph can be faster than coming up with all edges. Alternatively, given that LLMs obtain their best accuracy in pairwise causal discovery tasks, another way is use LLMs to critique human-generated graphs. LLMs can be used to iteratively critique the presence or absence of edges and help the analyst discover any missing edges.

Apart from graph creation, the other challenging task is to validate the output of any causal analysis, since the output depends on the initial causal assumptions (e.g., the conditional independence assumptions implied by the graph).

Here LLMs can act as useful assistants and help the analyst through chat conversation towards finding any flaws in the research and/or planning a robustness test. LLMs can also help in suggesting specific robustness checks for a causal analysis, such as the variables to use as negative controls. In Suppl. F.2, we present a case study of a conversation with an LLM to identify negative controls for an observational study of vaccine efficacy.

LLM + Causal tools: A new pipeline for causal analysis While LLMs can help with framing and validating causal assumptions—the steps in a causal analysis where human expertise is central—they are not capable of replacing the tools that do statistical analysis for estimating causal quantities. These tools (e.g., DoWhy, EconML, Ananke) implement principled algorithms from graphical causal models and the potential outcomes literature and are a workhorse for a causal analyst. We envision that LLMs can enable an end-to-end pipeline for causal inference where the code control seamlessly transfers between an LLM and causal tools. Practitioners can construct pipelines where graph creation is done in LLM-assisted interfaces while the final output graph is passed in to a statistical inference tool. Given a graph, an LLM may also be used to generate the code for conducting downstream analysis with a specified causal tool. Based on recent research (Schick et al., 2023), it may also be possible to augment LLMs and enable them to call out relevant causal tools and return the answer in a conversational chat interface. Suppl. F.1 presents a case study conversation with GPT-4 asking it to create sample code for synthetic data generation and causal analysis using the PyWhy open-source libraries.

LLMs as a fluid conversational interface, merging covariance- and logic-based reasoning. In practice when tackling complex causal challenges, people strategically alternate between different modes and methods for causal reasoning. People may begin with a logical reasoning approach to frame a high-level problem, then toggle to a covariance based approach to solve a subtask. They may switch to a mental simulation to consider counterfactuals, or validate consequences against logical assumptions. And they may continue interchanging complementary analytical methods repeatedly before honing in on their final answer. This kind of fluid iteration between modes of thinking to date required humans to translate their premises and knowledge across frameworks. Now, LLMs may be able to do more of that.

If so, this may allow a merging of the many kinds of causality into a single higher-level abstraction, allowing people to ask and have answered a much broader and

higher-level set of causal questions than could have been addressed through computational methods previously. A first step towards this goal can be to include the information about variables into the causal discovery problem.

5.2 New research questions on causality and LLMs

TL;DR: LLMs are enabling significant new advances in integrating domain knowledge and enabling natural, flexible interactions for causal tasks. Further research is needed to explore LLMs’ capabilities in causal discovery, effect inference, actual causality, human-LLM collaboration, and understanding and improving causal reasoning.

Our work is a proof of concept that LLMs hold the potential to provide domain knowledge and enable natural and flexible interactions for causal tasks. It will require further advances to enable effective human-AI interaction to realize this potential.

Knowledge-based causal discovery. We acknowledge that our formulation of knowledge-based pairwise discovery is different from the standard formulation. However, given that LLMs demonstrate non-trivial performance of inferring causal relationships based on variable names alone, we believe that this can have substantial implications on causal discovery research. Specifically, these results beg the question of how we can reformulate the problem of causal discovery to best leverage large language models and existing knowledge in the large amount of texts.

Typical problem formulation of causal discovery aims to identify the causal graph between a set of variables based on (observational or experimental) data associated with these variables. Metadata about variables are completely ignored. However, our results indicate that this may be a mistake since most datasets provide both data values and metadata that can be useful for inferring at least a part of the full graph. One promising direction is to consider the output of an LLM as a prior for a causal discovery algorithm. The relevant question is to what extent having an LLM-based prior would increase the accuracy of covariance-based discovery algorithms, where (Choi et al., 2022) present some promising results. Another, more provocative direction is to recognize that current discovery algorithms struggle on real-world datasets and we can aim for redefining the causal discovery problem such that it includes the meta information about variables and the existing knowledge encoded through LLMs (Zhiheng et al., 2022). Developing new discovery algorithms that

make use of both metadata and covariance in the data can be a fruitful research direction, wherein the LLM may be used as a prior, as a critic during learning, and/or as a post-processor.

In addition to building new algorithms, LLMs can also facilitate building benchmarks for covariance-based causal discovery. As illustrated from our results, causal discovery benchmarks are usually created by close collaboration with domain experts, with a strong assumption that domain experts know the groundtruth causal graphs. However, this assumption does not hold because 1) Possible edges grow superlinearly with the number of variables, so it is unlikely for a human expert to go over all possible edges reliably; 2) More importantly, domain experts may not be able to precisely provide the causal link between two variables given their knowledge. LLMs can help identify potentially missing edge and correct mislabeled edge, or at least point out relations that require additional checking.

LLM-guided effect inference. Identifying the correct adjustment set (e.g., using the backdoor criterion) is one of the most important problems for causal effect inference. However, the validity of the backdoor criterion depends on access to the true graph. LLMs’ capability for causal discovery can be useful for effect inference and opens up a research question on how to use LLMs to infer valid backdoor sets. In addition to discovering the full graph, techniques that utilize partial graph structure (Entner et al., 2013; Cheng et al., 2022) to identify backdoor sets may also benefit from LLMs’ knowledge-based discovery for inferring specific edges. Similarly, it may be interesting to see if LLMs can be used to suggest potential instrumental variables for a causal task, given metadata about observed variables.

A second promising direction is in utilizing LLM’s domain knowledge to help build robustness and validation checks for a given causal analysis. Typically, this involves an analyst selecting appropriate robustness checks given a problem and then reasoning over the background context to identify the correct variables to use for the check. Can LLMs be used for determining which robustness tests may be applicable for a given problem and which specific variables are most suited to be used in those tests?

Systematizing actual causality and attribution. The ability to support actual causal inference is one of the most potentially disruptive capabilities of LLMs. In domains ranging from law to intelligence analysis, analysts have to first summarize key elements of text data from multiple sources, then synthesize explanations about the degree to which events contributed to other events, i.e., explanations

of why and how things happened. While it is well known LLMs excel at the former task, using LLMs for actual causality presents a major opportunity to impact the latter. Using LLMs for actual causality also has potential impact to root cause analysis and credit assignment in domains ranging from engineering to reinforcement learning.

Formal causal models for inferring actual causality questions have struggled formalizing the many elements of common sense background knowledge that rely on when judging actual causality (Halpern, 2016; Icard et al., 2017; Knobe, 2003; Henne et al., 2017). We’ve argued that LLMs can work with these background concepts directly in natural language, sidestepping the challenge of shoehorning them into a formal model.

The BIG-Bench causal judgment dataset (Suzgun et al., 2022) is ideal for testing LLMs ability to answer practical actual causality questions. LLMs have performed less well on this benchmark relative to best human performance as compared to the rapid saturation of other benchmarks. However, answering these types of questions involves reasoning over commonsense background knowledge concepts, such as necessity, sufficiency, and normality. Our analysis shows that LLMs do better at answering questions about these basic elements than answering high-level actual causal judgment questions directly. A promising direction of research is developing methods that guide LLMs to use these actual causality primitives to answer higher level actual causal judgment questions, perhaps using formal actual causality theory as a guide.

Human-LLM collaboration. A recurring theme throughout the paper is the promise of human-AI collaboration in causal problems. There are interesting research questions on how best to facilitate such human-LLM collaboration for building graphs. The interaction will likely be iterative for the maximum impact: LLMs may suggest graph edges, take feedback from a human, and also can give feedback on a manually generated graph. Given that building a graph is a common impediment to doing causal analyses, an improvement in this task will have important consequences for the widespread adoption of causal data analysis.

Understanding and improving causal reasoning in LLMs. Why LLMs demonstrate such causal capabilities, as well as the limits of these capabilities is not yet well understood. According to Pearl’s causal hierarchy (Pearl & Mackenzie, 2018), LLMs, trained only on observational data, should not be able to answer interventional or counterfactual queries. Willig et al. (2023) offer a potential explanation by suggesting that LLMs sidestep the causal

hierarchy by learning *correlations* of causal facts from the training data, rather than learning to reason causally. The same explanation may apply to our results in counterfactual inference over text, which will require further study. At the same time, we see unpredictable failure modes, even in tasks where LLMs obtain high accuracy. In summary, understanding the nature of causal reasoning in LLMs and how to improve their robustness is a key future question.

6 Conclusion

Human domain knowledge has always been a central piece in causal analysis. In this paper, we studied the capabilities of large language models and found that LLMs can provide value by mimicking that domain knowledge, as a result of being trained on vast amounts of human-generated text. This mimicking is a result of a complex training process so it is not explainable and neither predictable: an LLM can fail on some queries while succeeding to provide causal reasoning in others. What is remarkable is how few times that such errors happen: our evaluation finds that on average, LLMs can outperform state-of-the-art causal algorithms in graph discovery and counterfactual inference, and can systematize nebulous concepts like necessity and sufficiency of cause by operating solely on natural language input.

From a research perspective, these results open up more questions than they answer, and we provide a list of research questions at the intersection of LLMs and causality. At the same time, due to the demonstrated capabilities of LLMs, we foresee a marked impact of LLMs on the practice of causal analysis. We outlined how LLMs can help reduce the burden of human expertise in tasks like graph discovery, effect inference, and attribution. Another contribution of LLMs is in bridging the divide between covariance- and logic-based causal analysis. By providing a flexible natural language interface for answering causal queries, LLMs can serve to unify these two branches of causal analysis and allow analyses that seamlessly straddle both to answer real-world causal questions.

References

- AML-BABEL. Aml-babel benchmark: Causal judgement. https://aml-babel.com/benchmark/datasets/bigbench-hard/causal_judgement, accessed 2023-04-25.
- Atchison, J. W. and Vincent, H. K. Obesity and low back pain: relationships and treatment. *Pain Management*, 2 (1):79–86, 2012.
- Bender, E. M., Gebru, T., McMillan-Major, A., and Shmitchell, S. On the dangers of stochastic parrots: Can language models be too big? In *Proceedings of the 2021 ACM conference on fairness, accountability, and transparency*, pp. 610–623, 2021.
- Bubeck, S., Chandrasekaran, V., Eldan, R., Gehrke, J., Horvitz, E., Kamar, E., Lee, P., Lee, Y. T., Li, Y., Lundberg, S., et al. Sparks of artificial general intelligence: Early experiments with gpt-4. *arXiv preprint arXiv:2303.12712*, 2023.
- Budhathoki, K., Janzing, D., Bloebaum, P., and Ng, H. Why did the distribution change? In Banerjee, A. and Fukumizu, K. (eds.), *Proceedings of The 24th International Conference on Artificial Intelligence and Statistics*, volume 130 of *Proceedings of Machine Learning Research*, pp. 1666–1674. PMLR, 13–15 Apr 2021.
- Byrne, R. M. *The Rational Imagination: How People Create Alternatives to Reality*. MIT Press, Cambridge, MA, 2005.
- Carlini, N., Ippolito, D., Jagielski, M., Lee, K., Tramer, F., and Zhang, C. Quantifying memorization across neural language models. *arXiv preprint arXiv:2202.07646*, 2022.
- Chang, W. Connecting counterfactual and physical causation. In *Proceedings of the Annual Meeting of the Cognitive Science Society*, volume 31, 2009.
- Chen, M., Tworek, J., Jun, H., Yuan, Q., de Oliveira Pinto, H. P., Kaplan, J., Edwards, H., Burda, Y., Joseph, N., Brockman, G., Ray, A., Puri, R., Krueger, G., Petrov, M., Khlaaf, H., Sastry, G., Mishkin, P., Chan, B., Gray, S., Ryder, N., Pavlov, M., Power, A., Kaiser, L., Bavarian, M., Winter, C., Tillet, P., Such, F. P., Cummings, D., Plappert, M., Chantzis, F., Barnes, E., Herbert-Voss, A., Guss, W. H., Nichol, A., Paino, A., Tezak, N., Tang, J., Babuschkin, I., Balaji, S., Jain, S., Saunders, W., Hesse, C., Carr, A. N., Leike, J., Achiam, J., Misra, V., Morikawa, E., Radford, A., Knight, M., Brundage, M., Murati, M., Mayer, K., Welinder, P., McGrew, B., Amodei, D., McCandlish, S., Sutskever, I., and Zaremba, W. Evaluating large language models trained on code, 2021.
- Cheng, D., Li, J., Liu, L., Yu, K., Le, T. D., and Liu, J. Toward unique and unbiased causal effect estimation from data with hidden variables. *IEEE Transactions on Neural Networks and Learning Systems*, 2022.

- Chockler, H. and Halpern, J. Y. Responsibility and blame: A structural-model approach. *Journal of Artificial Intelligence Research*, 22:93–115, 2004.
- Choi, K., Cundy, C., Srivastava, S., and Ermon, S. Lmpriors: Pre-trained language models as task-specific priors. *arXiv preprint arXiv:2210.12530*, 2022.
- Danilevsky, M., Qian, K., Aharonov, R., Katsis, Y., Kawas, B., and Sen, P. A survey of the state of explainable ai for natural language processing, 2020.
- Devlin, J., Chang, M.-W., Lee, K., and Toutanova, K. Bert: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pp. 4171–4186, 2019.
- Entner, D., Hoyer, P., and Spirtes, P. Data-driven covariate selection for nonparametric estimation of causal effects. In *Artificial intelligence and statistics*, pp. 256–264. PMLR, 2013.
- Frohberg, J. and Binder, F. CRASS: A novel data set and benchmark to test counterfactual reasoning of large language models. In *Proceedings of the Thirteenth Language Resources and Evaluation Conference*, pp. 2126–2140, Marseille, France, June 2022. European Language Resources Association. URL <https://aclanthology.org/2022.lrec-1.229>.
- Gerstenberg, T., Goodman, N., Lagnado, D., and Tenenbaum, J. From counterfactual simulation to causal judgment. In *Proceedings of the annual meeting of the cognitive science society*, volume 36, 2014.
- Ghazal, A., Ivanov, T., Kostamaa, P., Crolotte, A., Voong, R., Al-Kateb, M., Ghazal, W., and Zicari, R. V. Bigbench v2: the new and improved bigbench. In *2017 IEEE 33rd International Conference on Data Engineering (ICDE)*, pp. 1225–1236. IEEE, 2017.
- Glymour, C., Zhang, K., and Spirtes, P. Review of causal discovery methods based on graphical models. *Frontiers in genetics*, 10:524, 2019.
- Halpern, J. Y. *Actual causality*. MIT Press, 2016.
- Halpern, J. Y. and Hitchcock, C. Graded causation and defaults. *The British Journal for the Philosophy of Science*, 2015.
- Hegarty, M. Mechanical reasoning by mental simulation. *Trends in cognitive sciences*, 8(6):280–285, 2004.
- Hellner, J. Causality and causation in law. *Scandinavian studies in law*, 40:111–134, 2000.
- Henne, P. and O’Neill, K. Double prevention, causal judgments, and counterfactuals. *Cognitive science*, 46(5): e13127, 2022.
- Henne, P., Pinillos, Á., and De Brigard, F. Cause by omission and norm: Not watering plants. *Australasian Journal of Philosophy*, 95(2):270–283, 2017.
- Hernán, M. A. and Robins, J. M. Causal inference, 2010.
- Hitchcock, C. Portable causal dependence: A tale of concilience. *Philosophy of Science*, 79(5):942–951, 2012.
- Hobbhahn, M., Lieberum, T., and Seiler, D. Investigating causal understanding in llms. In *NeurIPS ML Safety Workshop*, 2022.
- Hoyer, P., Janzing, D., Mooij, J. M., Peters, J., and Schölkopf, B. Nonlinear causal discovery with additive noise models. *Advances in neural information processing systems*, 21, 2008.
- Huang, Y., Kleindessner, M., Munishkin, A., Varshney, D., Guo, P., and Wang, J. Benchmarking of data-driven causality discovery approaches in the interactions of arctic sea ice and atmosphere. *Frontiers in Big Data*, 4, 2021. ISSN 2624-909X. doi: 10.3389/fdata.2021.642182. URL <https://www.frontiersin.org/articles/10.3389/fdata.2021.642182>.
- Icard, T. and Goodman, N. D. A resource-rational approach to the causal frame problem. In *CogSci*, 2015.
- Icard, T. F., Kominsky, J. F., and Knobe, J. Normality and actual causal strength. *Cognition*, 161:80–93, 2017.
- Imbens, G. W. and Rubin, D. B. *Causal inference in statistics, social, and biomedical sciences*. Cambridge University Press, 2015.
- Jeannerod, M. Neural simulation of action: a unifying mechanism for motor cognition. *Neuroimage*, 14(1): S103–S109, 2001.
- Kahneman, D. and Miller, D. T. Norm theory: Comparing reality to its alternatives. *Psychological Review*, 93(2): 136–153, 1986.
- Kaiser, M. and Sipos, M. Unsuitability of notears for causal graph discovery when dealing with dimensional quantities. *Neural Processing Letters*, 54(3):1587–1595, 2022.

- Katz, D. M., Bommarito, M. J., Gao, S., and Arredondo, P. Gpt-4 passes the bar exam. *Available at SSRN 4389233*, 2023.
- Knobe, J. Intentional action and side effects in ordinary language. *Analysis*, 63(3):190–194, 2003.
- Knobe, J. and Shapiro, S. Proximate cause explained. *The University of Chicago Law Review*, 88(1):165–236, 2021.
- Kominsky, J. F., Phillips, J., Gerstenberg, T., Lagnado, D., and Knobe, J. Causal superseding. *Cognition*, 137: 196–209, 2015.
- Kosoy, E., Chan, D. M., Liu, A., Collins, J., Kaufmann, B., Huang, S. H., Hamrick, J. B., Canny, J., Ke, N. R., and Gopnik, A. Towards understanding how machines can learn causal overhypotheses. *arXiv preprint arXiv:2206.08353*, 2022.
- Kueffner, K. R. A comprehensive survey of the actual causality literature. 2021.
- Lee, P., Bubeck, S., and Petro, J. Benefits, limits, and risks of gpt-4 as an ai chatbot for medicine. *New England Journal of Medicine*, 388(13):1233–1239, 2023a.
- Lee, P., Goldberg, C., and Kohane, I. *The AI Revolution in Medicine: GPT-4 and Beyond*. Pearson, 1 edition, 2023b. ISBN ISBN-13: 9780138200138.
- Lipsitch, M., Tchetgen, E. T., and Cohen, T. Negative controls: a tool for detecting confounding and bias in observational studies. *Epidemiology (Cambridge, Mass.)*, 21(3):383, 2010.
- Liu, H., Ning, R., Teng, Z., Liu, J., Zhou, Q., and Zhang, Y. Evaluating the logical reasoning ability of chatgpt and gpt-4, 2023.
- Lombrozo, T. Causal-explanatory pluralism: How intentions, functions, and mechanisms influence causal ascriptions. *Cognitive psychology*, 61(4):303–332, 2010.
- Long, S., Schuster, T., Piché, A., Research, S., et al. Can large language models build causal graphs? *arXiv preprint arXiv:2303.05279*, 2023.
- Marcus, G. How come gpt can seem so brilliant one minute and so breathtakingly dumb the next?, 2022. URL <https://garymarcus.substack.com/p/how-come-gpt-can-seem-so-brilliant>.
- Marx, A. and Vreeken, J. Telling cause from effect using mdl-based local and global regression. In *2017 IEEE international conference on data mining (ICDM)*, pp. 307–316. IEEE, 2017.
- Mooij, J. M., Peters, J., Janzing, D., Zscheischler, J., and Schölkopf, B. Distinguishing cause from effect using observational data: Methods and benchmarks. *Journal of Machine Learning Research*, 17(32):1–102, 2016. URL <http://jmlr.org/papers/v17/14-518.html>.
- Moore, M. S. *Causation and responsibility: An essay in law, morals, and metaphysics*. Oxford University Press on Demand, 2009.
- Nadelhoffer, T. Desire, foresight, intentions, and intentional actions: Probing folk intuitions. *Journal of Cognition and Culture*, 6(1-2):133–157, 2006.
- Nauta, M., Bucur, D., and Seifert, C. Causal discovery with attention-based convolutional neural networks. *Machine Learning and Knowledge Extraction*, 1(1):312–340, 2019.
- Nguyen, N. and Nadi, S. An empirical evaluation of github copilot’s code suggestions. In *Proceedings of the 19th International Conference on Mining Software Repositories, MSR ’22*, pp. 1–5, New York, NY, USA, 2022. Association for Computing Machinery. ISBN 9781450393034. doi: 10.1145/3524842.3528470. URL <https://doi.org/10.1145/3524842.3528470>.
- Nori, H., King, N., McKinney, S. M., Carignan, D., and Horvitz, E. Capabilities of gpt-4 on medical challenge problems. *arXiv preprint arXiv:2303.13375*, 2023.
- Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C., Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A., et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35:27730–27744, 2022.
- Pearl, J. Causal inference in statistics: An overview. 2009a.
- Pearl, J. *Causality: Models, Reasoning, and Inference*. Cambridge University Press, 2 edition, 2009b.
- Pearl, J. *Causality*. Cambridge university press, 2009c.
- Pearl, J. and Mackenzie, D. *The book of why: the new science of cause and effect*. Basic books, 2018.
- Peters, J., Janzing, D., and Schölkopf, B. *Elements of causal inference: foundations and learning algorithms*. The MIT Press, 2017.

- Phillips, J., Luguri, J. B., and Knobe, J. Unifying morality's influence on non-moral judgments: The relevance of alternative possibilities. *Cognition*, 145:30–42, 2015.
- Sanh, V., Webson, A., Raffel, C., Bach, S. H., Sutawika, L., Alyafeai, Z., Chaffin, A., Stiegler, A., Scao, T. L., Raja, A., et al. Multitask prompted training enables zero-shot task generalization. *arXiv preprint arXiv:2110.08207*, 2021.
- Schick, T., Dwivedi-Yu, J., Dessì, R., Raileanu, R., Lomeli, M., Zettlemoyer, L., Cancedda, N., and Scialom, T. Toolformer: Language models can teach themselves to use tools. *arXiv preprint arXiv:2302.04761*, 2023.
- Sharma, A., Syrgkanis, V., Zhang, C., and Kıcıman, E. Dowhy: Addressing challenges in expressing and validating causal assumptions. *arXiv preprint arXiv:2108.13518*, 2021.
- Sharma, A., Li, H., and Jiao, J. The counterfactual-shapley value: Attributing change in system metrics. *arXiv preprint arXiv:2208.08399*, 2022.
- Shimizu, S., Hoyer, P. O., Hyvärinen, A., Kerminen, A., and Jordan, M. A linear non-gaussian acyclic model for causal discovery. *Journal of Machine Learning Research*, 7(10), 2006.
- Sinha, S., Chen, H., Sekhon, A., Ji, Y., and Qi, Y. Perturbing inputs for fragile interpretations in deep natural language processing, 2021.
- Sloman, S. A. and Lagnado, D. Causality in thought. *Annual review of psychology*, 66:223–247, 2015.
- Smith, G. T. On construct validity: issues of method and measurement. *Psychological assessment*, 17(4):396, 2005.
- Srivastava, A., Rastogi, A., Rao, A., Shoeb, A. A. M., Abid, A., Fisch, A., Brown, A. R., Santoro, A., Gupta, A., Garriga-Alonso, A., et al. Beyond the imitation game: Quantifying and extrapolating the capabilities of language models. *arXiv preprint arXiv:2206.04615*, 2022.
- Suzgun, M., Scales, N., Schärli, N., Gehrmann, S., Tay, Y., Chung, H. W., Chowdhery, A., Le, Q. V., Chi, E. H., Zhou, D., , and Wei, J. Challenging big-bench tasks and whether chain-of-thought can solve them. *arXiv preprint arXiv:2210.09261*, 2022.
- Tagasovska, N., Chavez-Demoulin, V., and Vatter, T. Distinguishing cause from effect using quantiles: Bivariate quantile causal discovery. In III, H. D. and Singh, A. (eds.), *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pp. 9311–9323. PMLR, 13–18 Jul 2020. URL <https://proceedings.mlr.press/v119/tagasovska20a.html>.
- Tomić, S., Soldo-Butković, S., Kovač, B., Faj, D., Jurić, S., Mišević, S., Knežević, L., and Vukašinović, D. Lumbar radiculopathy—factors effects on it's severity. *Collegium antropologicum*, 33(1):175–178, 2009.
- Tu, R., Zhang, K., Bertilson, B., Kjellstrom, H., and Zhang, C. Neuropathic pain diagnosis simulator for causal discovery algorithm evaluation. *Advances in Neural Information Processing Systems*, 32, 2019.
- Tu, R., Ma, C., and Zhang, C. Causal-discovery performance of chatgpt in the context of neuropathic pain diagnosis. *arXiv preprint arXiv:2301.13819*, 2023.
- Wei, J., Tay, Y., Bommasani, R., Raffel, C., Zoph, B., Borgeaud, S., Yogatama, D., Bosma, M., Zhou, D., Metzler, D., Chi, E. H., Hashimoto, T., Vinyals, O., Liang, P., Dean, J., and Fedus, W. Emergent abilities of large language models, 2022a.
- Wei, J., Wang, X., Schuurmans, D., Bosma, M., Chi, E., Le, Q., and Zhou, D. Chain of thought prompting elicits reasoning in large language models. *arXiv preprint arXiv:2201.11903*, 2022b.
- Willig, M., Zečević, M., Dhamsi, D. S., and Kersting, K. Can foundation models talk causality? *arXiv preprint arXiv:2206.10591*, 2022.
- Willig, M., ZEČEVIĆ, M., DHAMI, D. S., and KERSTING, K. Causal parrots: Large language models may talk causality but are not causal. 2023.
- Woodward, J. Sensitive and insensitive causation. *The Philosophical Review*, 115(1):1–50, 2006.
- Wu, P. and Fukumizu, K. Causal mosaic: Cause-effect inference via nonlinear ica and ensemble method. In *International Conference on Artificial Intelligence and Statistics*, pp. 1157–1167. PMLR, 2020.
- Yu, Y., Chen, J., Gao, T., and Yu, M. Dag-gnn: Dag structure learning with graph neural networks. In *International Conference on Machine Learning*, pp. 7154–7163. PMLR, 2019.
- Zhang, C., Bauer, S., Bennett, P., Gao, J., Gong, W., Hilmkil, A., Jennings, J., Ma, C., Minka, T., Pawlowski, N., and Vaughan, J. Understanding causality with large language models: Feasibility and opportunities, 2023.

Zhang, K. and Chan, L.-W. Extensions of ica for causality discovery in the hong kong stock market. In *Neural Information Processing: 13th International Conference, ICONIP 2006, Hong Kong, China, October 3-6, 2006. Proceedings, Part III* 13, pp. 400–409. Springer, 2006.

Zhang, K. and Hyvärinen, A. Causality discovery with additive disturbances: An information-theoretical perspective. In *Machine Learning and Knowledge Discovery in Databases: European Conference, ECML PKDD 2009, Bled, Slovenia, September 7-11, 2009, Proceedings, Part II* 20, pp. 570–585. Springer, 2009.

Zhang, K. and Hyvarinen, A. On the identifiability of the post-nonlinear causal model. *arXiv preprint arXiv:1205.2599*, 2012.

Zheng, X., Aragam, B., Ravikumar, P. K., and Xing, E. P. Dags with no tears: Continuous optimization for structure learning. *Advances in neural information processing systems*, 31, 2018.

Zhiheng, L., Jin, Z., Mihalcea, R., Sachan, M., and Schölkopf, B. Can large language models distinguish cause from effect? In *UAI 2022 Workshop on Causal Representation Learning*, 2022.

Zhong, W., Cui, R., Guo, Y., Liang, Y., Lu, S., Wang, Y., Saied, A., Chen, W., and Duan, N. Agieval: A human-centric benchmark for evaluating foundation models, 2023.

Zhou, Y., Muresanu, A. I., Han, Z., Paster, K., Pitis, S., Chan, H., and Ba, J. Large language models are human-level prompt engineers. *arXiv preprint arXiv:2211.01910*, 2022.

A Causal discovery: Additional details

A.1 Example prompts for the Tübingen benchmark

Two prompts per pair

Template:

- Does changing {A} cause a change in {B}? Please answer in a single word: yes or no.
- Does changing {B} cause a change in {A}? Please answer in a single word: yes or no.

Examples:

- Does changing the altitude cause a change in temperature? Please answer in a single word: yes or no.
- Does changing the temperature cause a change in altitude? Please answer in a single word: yes or no.

- Does changing a ball’s speed at the beginning of a ball track cause a change in its speed at the end of the track? Please answer in a single word: yes or no.

- Does changing a ball’s speed at the end of a ball track cause a change in its speed at the beginning of the track? Please answer in a single word: yes or no.
-

Single prompt

Template:

- Which cause-and-effect relationship is more likely?
- A. changing {A} causes a change in {B}.
- B. changing {B} causes a change in {A}.

Let’s work this out in a step by step way to be sure that we have the right answer. Then provide your final answer within the tags <Answer>A/B</Answer>.

Example:

- Which cause-and-effect relationship is more likely?
- A. changing the altitude causes a change in temperature.
- B. changing the temperature causes a change in altitude.

Let’s work this out in a step by step way to be sure that we have the right answer. Then provide your final answer within the tags <Answer>A/B</Answer>.

Table 14: Two kinds of prompt templates for the Tübingen benchmark. The first asks two questions per variable pair whereas the second (“Single prompt”) asks a single question to orient each pairwise edge.

A.2 Example LLM outputs for the Neuro-pathic pain dataset

An example of a correct output from gpt-3.5-turbo.

LLM: First, we need to understand the terms involved in the cause-and-effect relationship.

Left T6 Radiculopathy refers to a condition where there is damage or irritation to the nerve root that exits the spinal cord at the level of T6 on the left side of the body.

DLS T5-T6 refers to degenerative changes in the discs, ligaments, and/or bones at the level of T5-T6 in the spine.

Now, let's consider the two options:

A. Left T6 Radiculopathy causes DLS T5-T6. This would mean that the damage or irritation to the nerve root at T6 on the left side of the body somehow leads to degenerative changes in the discs, ligaments, and/or bones at the level of T5-T6. While it is possible that nerve damage could contribute to degenerative changes in the spine, it is more likely that the opposite is true.

B. DLS T5-T6 causes Left T6 Radiculopathy. This option suggests that the degenerative changes in the discs, ligaments, and/or bones at the level of T5-T6 are causing the Left T6 Radiculopathy. This is a more likely cause-and-effect relationship, as degenerative changes in the spine can lead to nerve root irritation or compression.

Therefore, the more likely cause-and-effect relationship is: <Answer>B</Answer>DLS T5-T6 causes Left T6 Radiculopathy.

An example of an incorrect output.

LLM: First, we need to understand what DLS T5-T6 and Right T6 Radiculopathy are. DLS stands for degenerative lumbar spondylosis, which is a condition where the spinal discs in the lumbar region of the spine degenerate over time. T5-T6 refers to the thoracic vertebrae in the middle of the spine. Radiculopathy is a condition where there is compression or irritation of a nerve root, which can cause pain, numbness, or weakness in the area of the body that the nerve supplies.

Based on this information, it is more likely that option B is the correct cause-and-effect relationship. This is because radiculopathy is typically caused by compression or irritation of a nerve root, which can then lead to pain or other symptoms in the area of the body that the nerve supplies. In this case, the Right T6 Radiculopathy is causing symptoms in the area of the body that corresponds to the DLS T5-T6 region of the spine.

Therefore, the correct answer is <Answer>B</Answer>.

B Evaluation: LLMs and counterfactual reasoning

B.1 Details on CRASS counterfactual evaluation

For computing accuracy, the LLM obtains a score of 1 if it outputs the correct answer option and 0 if it outputs an incorrect option. If the LLM output does contain an answer among the options A/B/C, we consider as if the LLM has skipped the question and assign a score according to the probability of obtaining the correct answer under random chance (0.33).

B.1.1 Example prompt

Here is an example prompt.

PROMPT: A woman sees a fire. What would have happened if the woman had touched the fire?

A: She would have seen fire.

B: That is not possible.

C: She would not have been burned.

D: She would have been burned.

Provide your answer within the tags, <Answer>A/B/C/D</Answer>.

B.1.2 Example responses from gpt-4

Here we list three example responses from gpt-4 that do not match the benchmark answer. In the first two examples, the answer from gpt-4 can be considered correct.

SYSTEM: You are a helpful assistant for counterfactual reasoning.

USER: A man catches a ball. What would have happened if the man had caught a water balloon?

A: He would have caught the ball and not the balloon.

B: He would have got wet.

C: He would have died.

D: He would have remained dry.

Let's work this out in a step by step way to be sure that we have the right answer. Then provide your final answer within the tags, <Answer>A/B/C/D</Answer>.

GPT-4: Step 1: Identify the difference between the two scenarios.

In the first scenario, the man catches a ball. In the second scenario, the man catches a water balloon.

Step 2: Determine the consequences of the difference. Catching a ball typically does not result in getting wet or any other significant consequences. However,

catching a water balloon may cause it to burst, potentially getting the man wet.

Step 3: Evaluate the possible outcomes.

A: He would have caught the ball and not the balloon. (This outcome does not address the counterfactual scenario where the man catches a water balloon.)

B: He would have got wet. (This outcome is possible if the water balloon bursts upon being caught.)

C: He would have died. (This outcome is highly unlikely and not a reasonable consequence of catching a water balloon.)

D: He would have remained dry. (This outcome is possible if the man catches the water balloon without it bursting.)

Step 4: Choose the most likely outcome.

The most likely outcome is that the man would have either gotten wet or remained dry, depending on whether the water balloon burst or not. Since we cannot determine with certainty whether the water balloon would burst, we cannot definitively choose between options B and D.

<Answer>Uncertain between B and D</Answer>

SYSTEM: You are a helpful assistant for counterfactual reasoning.

USER: A man does not get very nervous. What would have happened if he had gotten very nervous?

A: That is not possible.

B: The man would have passed out.

C: The man would not have passed out.

Let's work this out in a step by step way to be sure that we have the right answer. Then provide your final answer within the tags, <Answer>A/B/C</Answer>.

GPT-4: Step 1: Identify the counterfactual situation - The man gets very nervous.

Step 2: Consider the possible outcomes of the counterfactual situation.

a) The man passes out due to extreme nervousness.

b) The man does not pass out but experiences other symptoms of nervousness.

Step 3: Evaluate the likelihood of each outcome.

a) Passing out due to extreme nervousness is possible, but not guaranteed.

b) It is more likely that the man would experience other symptoms of nervousness without necessarily passing out.

Final Answer: <Answer>C</Answer>

Below we show an example where people may easily understand that the man is walking on his bed indoors, but the LLM incorrectly assumes that he continues walking *outdoors* on a bed towards the same destination.

SYSTEM: You are a helpful assistant for counterfactual reasoning.

USER: A man walks on a street. What would have happened if a man had walked on a bed?

A: He would have been late.

B: Nothing special would have happened.

C: He would have arrived on time.

Let's work this out in a step by step way to be sure that we have the right answer. Then provide your final answer within the tags, <Answer>A/B/C</Answer>.

GPT-4: Step 1: Consider the difference between walking on a street and walking on a bed.

- Walking on a street is a normal activity and allows for efficient movement from one place to another.

- Walking on a bed is unusual and would likely slow down the person's movement due to the soft and unstable surface.

Step 2: Evaluate each option based on this difference.

A: He would have been late.

- This option suggests that walking on a bed would have slowed the man down, making him late for his destination.

B: Nothing special would have happened.

- This option implies that walking on a bed would not have had any significant impact on the man's journey.

C: He would have arrived on time.

- This option suggests that walking on a bed would not have affected the man's arrival time at his destination.

Step 3: Determine the most likely outcome.

- Given that walking on a bed is an unusual and inefficient way to move, it is most likely that the man would have been slowed down and potentially arrived late at his destination.

<Answer>A</Answer>

C Evaluation: Necessary and Sufficient Causes

C.1 Classes of actual causality vignette scenarios

For a discussion of the different scenario classes, see Chapter 3 in Kueffner (2021). Below we provide a brief summary.

1. Symmetric Overdetermination: Multiple processes that lead to the same outcome, complete at the same

time.

2. Switch: An action chooses between multiple processes, each of which lead to the same outcome. Thus, the action becomes immaterial for the outcome.
3. Late preemption: Two processes are running in parallel. Both would produce the same outcome, but one completes before the other.
4. Early preemption: Two processes would produce the same outcome, but one completes before the other has started.
5. Double preemption: A focal process leads to an event. An action that would have stopped the focal process, is itself stopped by another process.
6. Bogus preemption: An action stops a focal process, but the focal process itself is inactive wrt its effect on the outcome.
7. Short circuit: An original action makes the focal process inactive. An action is taken to prevent this inactive process; it triggers the process but the process is unable to complete because of the original action.

C.2 Prompt generation for necessary and sufficient causes

Given the complicated definitions of actual causality (see, for example, definitions in Chapter 2 of Halpern (2016)), it is non-trivial to express them in natural language. We thus turned to the LLM itself to suggest prompts to use.

Below we describe a method to generate an appropriate prompt for a given task, based on Zhou et al. (2022). For generating a prompt that instructs an LLM to identify the necessary cause, we provide a meta-prompt that gives a few demonstrations of the task (vignettes as input and the necessary cause as output) and then asks the LLM to generate the instruction for that task. We then use a paraphrased version of the LLM’s instruction to create our prompt.

Specifically, we provide three example vignettes with the scenario and correct actual cause and ask the LLM to output the instructions that must have led to the output in each example. After a few executions, the LLM provided a reasonable principle called “minimal rewriting”. We tweaked the wording to the more common “minimal change” principle and use it in our prompt.

Here is the meta-prompt and its response from gpt-3.5-turbo.

META-PROMPT: I taught my friend some concepts in counterfactual reasoning. Based on my teaching, they provided the following output, given an input scenario and candidate outputs.

Input: A window shatters when it is hit by a bullet. Alice and Bob each fire a bullet at a window, simultaneously striking the window, shattering it.

Candidate outputs: Both Alice and Bob, Any of Alice or Bob

Output: Both Alice and Bob

Input: Alice, Bob, Carol, Dave and Eve all fire at a window. The window shatters after three hits.

Candidate outputs: All five individuals, any individual, the first three individuals who fired

Output: The first three individuals who fired

Input: Alice fills a sink with water. At each time interval Alice adds another drop of water. At one point the sink overflows.

Candidate outputs: All droplets added to sink, Only the last droplet added to sink

Output: Only the last droplet added to sink

What logic was the friend using to select the correct output?

GPT-3.5-TURBO: : *It seems that your friend was applying the principle of minimal rewriting to select the correct output. The principle of minimal rewriting is a criterion used in counterfactual reasoning to evaluate which changes to the initial situation are necessary to obtain a given outcome.*

In this case, your friend was likely evaluating which minimal changes to the input scenario would lead to the candidate outputs. By comparing the input scenario to the candidate outputs, your friend identified which changes were necessary to obtain the desired outcome, and selected the output that required the least amount of change.

For example, in the first scenario, your friend identified that both Alice and Bob firing their bullets simultaneously was the minimal change needed to shatter the window. Therefore, they selected the output “Alice and Bob” as the correct answer, as it required the least amount of rewriting.

Similarly, in the second scenario, your friend identified that the minimal change to shatter the window was the first three individuals firing. Therefore, they selected “The first three individuals who fired” as the correct output.

Overall, your friends approach demonstrates a good understanding of counterfactual reasoning and the principle of minimal rewriting.

We follow the same procedure for creating a prompt for identifying the sufficient cause. One of the LLM outputs mentions the “multiple sufficient causes” principle, which we use in our prompt.

C.3 “Lab-vignettes”: A novel actual causality benchmark

To avoid LLM memorization issues, we construct a new dataset for inferring necessary and sufficient causes, based on the available dataset. Specifically, we perturb the details in the original dataset and adapt it for scenarios in a chemistry lab, involving reagents, mixtures and crystals. Table 15 shows example vignettes from this dataset. We apply our memorization test to validate that our newly created dataset cannot be recreated by the LLM.

C.4 Example responses from LLMs

C.4.1 Correct response on an Early Preemption vignette

We show the answers from gpt-3.5-turbo and GPT-4 on the example prompt in main text.

For necessary cause, both models yield the correct answer.

GPT-3.5-TURBO: : If Alice had not caught the passing cricket ball, the ball would have hit the solid brick wall and stopped there. Therefore, the window would not have been hit by the ball and would have remained intact. Hence, Alice is not a necessary cause for the window being intact.

<Answer>No</Answer>

GPT-4: : To determine whether Alice is a necessary cause for the window being intact, we can apply the principle of minimal change. Let’s consider the counterfactual situation where Alice does not catch the ball.

In the counterfactual situation, the ball would continue on its trajectory and hit the solid brick wall. As the brick wall would have stopped the ball, it would not have continued towards the window. Therefore, the window would still be intact even if Alice had not caught the ball.

<Answer>No</Answer>

For sufficiency, however, only GPT-4 provides the correct answer (Alice’s action is sufficient).

GPT-3.5-TURBO: : No, Alice is not a sufficient cause for the window being intact. While Alice catching the ball prevented it from hitting the brick wall, there could have been other factors that prevented the ball from breaking the window, such as the ball losing momentum or the window being made of sturdy material. Therefore, there are multiple sufficient causes for the window being intact. <Answer>No</Answer>.

GPT-4: : Using the principle of multiple sufficient causes, we can reason that there are multiple potential causes

that could result in the window remaining intact. In this scenario, Alice catching the ball is one sufficient cause, as it prevents the ball from hitting the window. Another sufficient cause would be the solid brick wall, which would have stopped the ball if Alice hadn’t caught it.

Given that there are multiple sufficient causes, Alice’s action of catching the ball is not the sole determinant of the window remaining intact. However, her action is one of the sufficient causes that contributed to the outcome.

<Answer>Yes</Answer>

C.4.2 Unpredictable failure mode: An example

Using the short circuit vignette, we show an example of gpt-4’s unpredictable failure modes. It applies the correct reasoning for most prompts, but applies an incorrect reasoning (assessing necessity only instead of sufficiency) for the following prompt. We show the input prompt and gpt-4’s answer below.

PROMPT: You are an expert in counterfactual reasoning. Given an event, use the principle of minimal change to answer the following question.

Input: Carol is alive (CA). Alice puts a harmless antidote in Carol’s water (AA). Adding antidote to the water, protects it against poison (WS - ‘water save’). If Alice puts the antidote into Carol’s water, Bob will poison the water (BP) Adding poison to an unprotected water makes it toxic (WT). If Carol would drink toxic water she would die (i.e. inhibiting CS). Carol consumes her water and survives (CS).

Is Alice a necessary cause for Carol’s survival?

After your reasoning, provide the final answer within the tags <Answer>Yes/No</Answer>.

GPT-4: :

“We can set up the causal chain as follows:

1. Alice puts the antidote in Carol’s water (AA).
2. Antidote in the water protects it against poison (WS).
3. If Alice puts the antidote into Carol’s water, Bob will poison the water (BP).
4. Adding poison to unprotected water makes it toxic (WT).
5. If Carol drinks toxic water, she would die (not CS).

To determine if Alice is a sufficient cause for Carol’s survival, let’s consider a counterfactual scenario:

If Alice had not put the antidote in Carol’s water, would Carol have survived?

In this scenario, we have:

1. Alice does not put the antidote in Carol’s water (not AA).
2. Antidote in the water does not protect it against poison (not WS).
3. Bob does not poison the water (not BP).
4. The water remains nontoxic (not WT).
5. Carol consumes her water and survives (CS).

Vignette Type	Input Context	Event	Actor	Nec.	Suff.
Overdetermination	There is a fire in the chemistry lab. A can of water would douse the fire. Agents X and Y both spray a can of water each, dousing the fire.	fire being doused	Agent X	No	Yes
Switch	Reagent X is added to a mixture, which leads to an explosion and kills Sam. Otherwise, Reagent Y in Sam's pocket would have infected him and killed him as well.	Sam's death	Reagent X	No	Yes
Late preemption	Any of Reagent X or Reagent Y can be added to a mixture to convert it into a crystal. Reagent X is added first and the mixture turns to crystal. Reagent Y is added later and but does not mix since the crystal is already formed.	crystal formation	Reagent X	No	Yes
Early preemption	Sam reaches out and catches the falling test tube containing a contaminating reagent. The next thing on the test tube's trajectory was a soft foam mattress which would have prevented the test tube from reaching the floor. Beyond that there was the floor.	floor being contaminated	Sam	No	Yes
Double preemption	Sam intends to add reagent X to a mixture that would convert it to a crystal. Riya intends to prevent Sam from adding the reagent. Riya tries to stop Sam from adding the reagent, but Riya is stopped by Frank. Sam adds the reagent and the crystal is formed.	crystal formation	Sam	Yes	No
Bogus preemption	Sam intends to put Reagent X in a mixture to make it explode. However, Sam does not put the reagent in the mixture. Bob puts an anti-explosion Reagent Y into the mixture. The mixture would explode if Reagent X is added without the addition of Reagent Y. Mixture does not explode.	mixture not exploding	Sam	No	Yes
Short circuit	A test tube contains a mixture. Sam puts sand in the mixture. Adding sand to the mixture protects it against explosion. Only if Sam puts sand in the mixture, Riya will add Reagent Y to the mixture. Adding Reagent Y to a mixture makes it explode, unless the mixture already contains sand. There is no explosion.	avoiding explosion	Sam	No	Yes
Miscellaneous	Sam is heating Reagent X in a test tube. If the heat is not stopped at the right time, the test tube would explode. Sam does not stop the heat. The test tube explodes.	explosion	Sam's inaction	Yes	Yes

Table 15: “lab-vignettes”: Examples of novel vignettes for evaluation of inferring necessary and sufficient causes. Each vignette is associated with two questions: “*Is {Actor} a necessary cause of {Event}?*” and “*Is {Actor} a sufficient cause of {Event}?*”

In this counterfactual scenario, Carol still survives even if Alice does not put the antidote in her water. This indicates that Alice’s action was not a sufficient cause for Carol’s survival, as Carol would have survived even without Alice’s intervention.

<Answer>No</Answer>

D Assessing Responsibility

So far we looked at necessary or sufficient actual causes, which are binary: an event is either a cause or not a cause of another event. In natural language, the causal judgments we make are relative, e.g., we might ascribe *more* causality to “smoking in the woods” than to “presence of dry leaves in Summer.” Formal actual causality models capture this with *graded causation*, nonbinary measures of actual causation that enable comparisons between causal events Halpern & Hitchcock (2015). In natural language we might use the simple concept of “responsibility.” Responsibility is a natural way to rank causal events in everyday language, e.g., “smoking in the woods was more responsible for the forest fire than the presence of dry leaves in Summer.”

Halpern (2016); Chockler & Halpern (2004) provide an SCM-based definition of responsibility; in simple terms, they define responsibility of event C for event E is $1/(N+1)$, where N is the minimal number of other causal events contributing to E that have had not happen to make C a necessary cause of E. But mapping the SCM definition of responsibility back to plain language is nontrivial, just as it was with model-based $P(E_{C=1} = 1 | C = 0, E = 0)$ -definition of sufficiency back to plain language was challenge. This SCM definition is purely a measure of causal contribution. But in everyday language, “responsibility” combines a degree of causal contribution with the concepts of intention, epistemic state, moral obligation, etc. mentioned earlier. So for this case study, we attempted to scope the definition of responsibility to focus on causal contribution.

Responsibility: The relative degree to which an event causally contributes to a particular outcome event, relative to other events that caused that outcome. Given two causal events that cause the same outcome event, one causal event can be more, less, or equally responsible relative to the other. This definition of responsibility does not depend on whether the causal event is a person’s action.

D.0.1 Inferring responsibility: Motivating Example

Taking our cue from the LLM’s ability to define relevant variables on the fly, our definition allows the LLM to provide a relative ranking of causes, rather than quantify an

absolute responsibility value. Further, we make it clear that responsibility is the same regardless of whether the causal event was a person’s action; we hope that if the definition of responsibility is the same for both inanimate object and people, then we cue the LLM to focus on causal contribution. Moreover, we cue the definition to ignore intention, morality, epistemic state, etc., or at the very least to limit those factors to levels that people usually ascribe inanimate objects when they are anthropomorphized in everyday language (e.g. “cheated by fate”, “the universe is telling you something”, “my phone is acting up,” etc.) We then prompted the LLM to use this definition for the beer spilling example.

PROMPT: In the context of this dialog, define “responsibility” as the relative degree to which an event causally contributes to a particular outcome event, relative to other events that caused that outcome. Given two causal events that cause the same outcome event, one causal event can be more, less, or equally responsible relative to the other. This definition of responsibility does not depend on whether the causal event is a person’s action.

Based on this definition of responsibility, make a one sentence statement about how responsible is Mike’s bumping of the table is for the spilling of the beer.

LLM: Based on this definition of responsibility, Mike’s bumping of the table is highly responsible for the spilling of the beer, as it was the immediate and direct cause that led to the beer bottle falling and spilling.

We now introduce two challenging scenarios. These are based on overdetermination and double preemption, two of the scenario types we saw in Section 4.2 for deciding necessary and sufficient causes.

Overdetermination. Overdetermination is one example where a simple counterfactual dependence definition of actuality falls short relative to human judgments. To illustrate we introduce a new information to the beer spilling story, that Susan bumped into the table shortly after Mike and with enough force to spill the beer had it still been there. With this additional information, Mike’s bump is no longer a necessary cause of the spill; if he hadn’t bumped the table, the beer still would have spilled. Despite the new lack of counterfactual dependence, humans still judge that Mike caused the spill. Analogously, if two children throw a rock at a window, most people judge that they both children are window breakers, regardless of whose rock struck first and broke the glass.

First we make sure the LLM can assess that counterfactual dependence is no longer present.

PROMPT: You said “if Mike hadn’t bumped into the table, the beer bottle would not have fallen and spilled.” Call this statement A. Suppose that Susan tripped and fell onto the table in the instant after Mike bumped into the table and the bottle started falling. If the bottle had not already fallen, Susan’s fall would have caused the bottle to fall. Assuming scenario A was true before, can statement A still be true? Provide a one sentence explanation why or why not.

LLM: No, scenario A can no longer be considered true because if Susan fell onto the table in the instant after Mike bumped it, Susan’s fall would have been the determining factor causing the beer bottle to fall and spill, regardless of Mike’s initial bump.

This result aligns with our intuition. Now, we expect that, given that there are two causes of the beer spilling, the LLM should be able downgrade Mike’s level of responsibility relative to before.

PROMPT: How, if at all, does this new information change your assessment of the responsibility of Mike’s table bump?

LLM: With this new information, the responsibility of Mike’s table bump for the spilling of the beer is reduced, as Susan’s fall would have also led to the same outcome, making both events contributing causal factors to the beer spill.

Double prevention. *Double prevention* describes another case where humans ascribe causality to events that are counterfactually independent of the outcome. In double prevention, we have an initial causal event C that causes the outcome. As before C is an event that causes outcome E. Now we introduce a new causal event chain $A \rightarrow B \rightarrow E$. If B occurs E does not occur. If A occurs B does not occur. Thus A prevents B from preventing E.

In our beer spilling example, Mike bumps the table (event C), which causes the beer bottle to fall (event E). Jack sees the bottle falling and tries to catch it (event B) and thus prevent event E. but Peter bumps into Jack (event A), which causes Jack to miss the bottle .

In this case, both A and E have counterfactual dependence; the outcome would not have occurred if A had not occurred. But humans tend to judge that Mike’s bumping the table caused the beer to spill and that Peter’s bumping into Jack, the double preventer, did not Chang (2009); Lombrozo (2010); Moore (2009). In the following prompt, we extend the beer spilling scene to include a double prevention scenario. In this case, Mike is bumps the table (event A), Jack tried to prevent it (event B), but Peter interferes with Jack’s attempt (event C). We expect that the LLM should assign more responsibility to Mike than Peter.

Dataset	% Memorized			
	GPT-3.5		GPT-4	
	Cells	Rows	Cells	Rows (%)
Tübingen	58.9%	19.8%	61%	25%
Lab-vignettes	–	–	–	0%
CRASS	–	–	–	–

Table 16: Our memorization test shows the Tübingen dataset is in the training dataset and has been at least partially memorized by GPT-3 and GPT-4; our experiments with the novel lab-vignettes show that the dataset has not been memorized. Results for CRASS and other benchmark datasets are forthcoming.

PROMPT: Suppose that Jack saw that the bottle was about to fall, and he reached out to catch the bottle. But just as Jack was about to catch the bottle, Peter bumped into Jack, which caused him to miss the bottle, which then fell to the ground. Who is more responsible for the bottle falling to the ground? Mike or Peter?

LLM: In this scenario, Peter is more responsible for the bottle falling to the ground, as his bumping into Jack directly interfered with Jack’s attempt to catch the bottle and prevent it from falling. While Mike’s bumping the table initially set the events in motion, Peter’s action had the most immediate impact on the outcome.

Here, we observe a failure of the ChatGPT4 to capture the nuance of double prevention at least by human standards.

E Probing LLM Behaviors Further

E.1 Memorization Tests

The system message used for the memorization test is:

SYSTEM: You are an AI assistant that has read many sources of text from the internet. I am looking at text from the <DATASET NAME AND DESCRIPTION>dataset, published by <AUTHORS>as <PUBLICATION INFO>. Here is the README for the dataset:
<README>

I am going to list some sample rows of data, and I want you to complete each row as best as possible. I am testing your memory.

USER: pair0005 Age

ASSISTANT: pair0005 Age Length Abalone
->

Session 1:

Q: What is $8 * 8 + 5 * 12$?

A: 104

Q: Please show your work

A: $8 * 8 = 64$
 $5 * 12 = 60$
 $64 + 60 = 104$

Session 2:

Q: What is $8 * 8 + 5 * 12$?

A: 110

Q: Please show your work

A: $8 * 8 = 64$
 $5 * 12 = 60$
 $64 + 60 = 110$

Figure 5: LLM completions are highlighted in yellow. Session 1: In a multi-turn interaction, GPT3.5 first gives an erroneous response to a math problem. Prompted to show its work, it correctly solves the sub-parts of the problem, but again gives the wrong final answer. Session 2: we probe for the influence of the first wrong answer on the final wrong answer by replaying an *interventional*-conversation and asking GPT3.5 to complete only the final answer.

E.2 Perturbation Test Example

We give a brief illustration of a perturbation experiment in Figure 5, which shows a GPT 3.5 model giving an erroneous answer to an arithmetic problem, even when correctly solving subparts of the problem when prompted. We might wonder whether the LLM has learned a pattern that the final answer must match its initial answer, instead of a more correct pattern that the final answer must match the preceding arithmetic expression. With LLMs, we can run a perturbation experiment to test this hypothesis, intervening on the first answer to study its influence on the final answer. By redacting then perturbing each of the tokens in the prompt, one by one, we can determine which tokens are leading to the error and better understand reliabilities. Once we have a target token or word, we can perturb its value systematically to identify how it influences the final answer. Table 17 shows that the perturbation to the first answer of Figure 5 is often directly reflected in the final result. Interestingly, as the perturbation value increases, the final answer is more likely to snap to the correct result instead.

F Supplementary Discussion

F.1 Example of LLM-assisted Causal Analysis Code Generation

Here, we present, without comment, a brief case study asking GPT-4 to generate python code for a demonstration causal analysis.

USER: could you give me a jupyter notebook that creates synthetic data for a causal problem (please make a simple DAG and data generator); then calls CMU’s causal-learn package to learn a dag; then calls DoWhy and EconML to estimate a specific treatment effect using the learned graph, and validates with sensitivity analyses

GPT-4: I cannot directly provide you with a Jupyter Notebook file, but I can give you a step-by-step guide on how to create one yourself. In this example, we’ll create a simple synthetic dataset with 3 variables, where $X1 \rightarrow X2 \rightarrow Y$ and $X1 \rightarrow Y$, and use causal-learn, DoWhy, and EconML to estimate the treatment effect of $X1$ on Y .

1. Install the required packages:

```
!pip install numpy pandas causal-learn dowhy econml
```

2. Import the necessary libraries:

```
import numpy as np
import pandas as pd
from causal_learn import causal_graph
from dowhy import CausalModel
from econml.metalearners import TLearner
from sklearn.linear_model import LinearRegression
```

3. Generate synthetic data:

```
np.random.seed(42)
n_samples = 1000

# Generating X1 (treatment variable)
X1 = np.random.normal(0, 1, n_samples)

# Generating X2 (mediator variable)
X2 = 0.5 * X1 + np.random.normal(0, 1, n_samples)

# Generating Y (outcome variable)
```

```

11 Y = 2 * X1 + 0.8 * X2 + np.
    random.normal(0, 1,
    n_samples)
12
13 # Create a dataframe
14 data = pd.DataFrame({'X1': X1, '
    X2': X2, 'Y': Y})

```

4. Learn the causal graph using the causal-learn package:

```

1 graph = causal_graph.CausalGraph
    (df=data)
2 graph.estimate_graph()
3 print(graph.graph_)

```

5. Use DoWhy to create a causal model:

```

1 # Create a causal model using
    the learned graph
2 causal_model = CausalModel(
3     data=data,
4     treatment='X1',
5     outcome='Y',
6     graph=graph.graph_.to_string
    ())
7 )
8
9 # Identify causal effect
10 identified_estimand =
    causal_model.identify_effect
    ()

```

6. Estimate the treatment effect using EconML:

```

1 # Train a T-learner using linear
    regression
2 t_learner = Tlearner(
    LinearRegression(),
    LinearRegression())
3 t_learner.fit(Y=data['Y'], T=
    data['X1'], X=data[['X2']])
4
5 # Estimate the treatment effect
6 treatment_effect = t_learner.
    const_marginal_effect(data[['
    X2']])

```

7. Validate the results using sensitivity analyses:

```

1 # Perform sensitivity analysis
    using DoWhy
2 refutation_results =
    causal_model.refute_estimate
    (
3     identified_estimand,
4     treatment_effect,
5     method_name='
        random_common_cause'

```

```

6 )
7
8 print(refutation_results)

```

This example demonstrates how to create synthetic data for a causal problem, learn a DAG using the causal-learn package, estimate treatment effects using DoWhy and EconML, and perform sensitivity analyses. Note that the actual treatment effect and learned treatment effect may differ due to noise in the data and the assumptions made by the methods used.

F.2 Example of LLM-assisted Identification of Negative Controls

Here, we present, without comment, a brief case study asking GPT-4 to generate negative and positive controls. We add list formatting for readability.

USER: In a causal observational study of vaccine efficacy, we are matching people who have voluntarily gotten a vaccine with individuals who have not yet done so. We are matching on attributes intended to capture all confounding effects. To help validate our analysis, what are some examples of negative controls where we might expect to see zero treatment effect?

GPT-4: In a causal observational study, negative controls are variables or outcomes for which you do not expect to see a treatment effect. By including negative controls in your analysis, you can help validate your matching process and assess the presence of unmeasured confounding. Here are some examples of negative controls where you might expect to see zero treatment effect in a study of vaccine efficacy:

Hair color: The vaccine should have no impact on the natural hair color of individuals, so you should expect no difference in hair color distribution between the vaccinated and unvaccinated groups.

Blood type: There should be no relationship between receiving a vaccine and an individual's blood type. So, the distribution of blood types should be similar in both groups.

Handedness: The vaccine should not affect whether a person is right-handed, left-handed, or ambidextrous, so you should expect no difference in handedness between the vaccinated and unvaccinated groups.

Taste preferences: Vaccination should not influence an individual's taste preferences for certain foods or flavors, so no difference should be observed in this regard between the two groups.

Birth month: There should be no relationship between receiving a vaccine and the month a person was born. Therefore, the distribution of birth months should be similar for vaccinated and unvaccinated individuals.

Past unrelated medical procedures: Previous medical procedures unrelated to vaccination, such as dental work or orthopedic surgeries, should not be affected by vaccination status. As such, there should be no difference in the prevalence of these procedures between the groups.

Favorite color: An individual's favorite color should not be influenced by vaccination status. Consequently, the distribution of favorite colors should be similar for both groups.

Including these negative controls in your analysis can help to evaluate whether your matching process is successful in capturing all confounding effects. If you observe a treatment effect in these negative controls, it may indicate that there are unmeasured confounding factors or issues with your matching process.

USER: In what subpopulations would we expect to see zero treatment effect on disease prevention? (for the disease being vaccinated against)

GPT-4: There are certain subpopulations where you might expect to see zero or limited treatment effect on disease prevention for the disease being vaccinated against:

Individuals with pre-existing immunity: People who have already been infected with the disease and developed immunity might not show additional benefits from vaccination. However, the vaccine may still help to boost their immunity or protect against different strains of the disease.

Individuals with specific immune deficiencies: Some people have immune system disorders or are taking immunosuppressive medications that may prevent them from mounting an effective immune response to the vaccine. In such cases, the vaccine may not provide any additional protection against the disease.

Non-responders to the vaccine: A small proportion of individuals might not develop an adequate immune response to the vaccine due to genetic factors or other unknown reasons. These non-responders may not experience any protection from the vaccine.

Individuals vaccinated after infection: If someone receives the vaccine after already being infected with the disease, the vaccine may not have a chance to provide any additional protection. This is because the immune response generated by the vaccine might not have enough time to develop before the disease progresses.

Individuals vaccinated with expired or improperly stored vaccines: If a vaccine has expired or has been stored improperly, it may lose its potency and not induce the desired immune response. In such cases, individuals receiving these vaccines may not experience any protection against the disease.

It's important to note that even if a subpopulation does not experience a treatment effect from the vaccine, the vaccine can still be beneficial for the overall

population. Widespread vaccination can lead to herd immunity, which can indirectly protect individuals in these subpopulations by reducing the overall prevalence of the disease.

USER: Are there any subpopulations in which we would expect a known but non-zero effect? to help validate our observational causal analysis?

GPT-4: Yes, there are subpopulations where you might expect a known but non-zero effect of vaccination on disease prevention. These groups can serve as positive controls to help validate your observational causal analysis. Some examples include:

Age groups: Vaccine efficacy can vary by age group. For instance, older adults or very young children may have a less robust immune response to vaccination compared to young adults. You may observe a reduced but non-zero effect in these age groups compared to others.

Individuals with certain chronic conditions: People with specific chronic medical conditions like diabetes, cardiovascular diseases, or respiratory illnesses might have a weaker immune response to vaccines. As a result, you may see a lower but non-zero vaccine effect in these subpopulations.

Pregnant individuals: Pregnancy can cause temporary changes in the immune system, which might lead to a reduced response to vaccines. Consequently, you may observe a smaller but non-zero effect in pregnant individuals compared to non-pregnant individuals.

Smokers: Smoking is known to have adverse effects on the immune system, potentially leading to a reduced immune response to vaccines. In this subpopulation, you may find a lower but non-zero vaccine effect compared to non-smokers.

Individuals receiving a lower vaccine dose or a different vaccination schedule: Some subpopulations might receive a lower vaccine dose or a different vaccination schedule, which can result in varying levels of protection. You may see a reduced but non-zero effect in these cases.

By comparing the treatment effect in these subpopulations to the overall population or other subpopulations with expected higher vaccine efficacy, you can help validate your causal analysis. If your analysis can correctly capture these known differences in treatment effects among various subpopulations, it provides evidence that your methodology is robust and can potentially account for confounding factors.

USER: if we have multiple observations per person over time. I.e., we know whether they are infected with the given disease on any given day after vaccination, then does that suggest any time-bound negative or positive controls?

GPT-4: Yes, when you have multiple observations per person over time (longitudinal data), you can use time-bound negative or positive controls to help validate

your causal analysis. Here are some examples:

Time-bound Negative Controls:

Pre-vaccination period: Compare the infection rates before vaccination between the future vaccinated and unvaccinated groups. Since neither group has received the vaccine at this stage, you should expect no difference in infection rates due to vaccination.

Short time window post-vaccination: Analyze the infection rates in a very short time window (e.g., within a few days) after vaccination. The vaccine's protective effect usually takes some time to develop, so you should expect minimal to no treatment effect in this short period.

Time-bound Positive Controls:

Peak immune response period: Identify the time window in which the vaccine's immune response is expected to be at its peak based on clinical trial data or previous research. Compare the infection rates between the vaccinated and unvaccinated groups during this period, and you should expect to see a substantial treatment effect.

Waning immunity period: Investigate the infection rates during a period when the vaccine's protective effect is known to wane or decrease. During this time, you might still expect to see a treatment effect, but it could be smaller compared to the peak immune response period.

Post-booster vaccination: If booster shots are administered to maintain or enhance immunity, you can compare infection rates between vaccinated individuals who have received boosters and those who have not (or unvaccinated individuals). You should expect to see a positive treatment effect in those who received the booster shots.

Using these time-bound controls can help you assess the validity of your causal analysis, as you expect specific patterns in the treatment effect over time. If your analysis can capture these expected temporal patterns, it provides evidence that your methodology is reliable and able to account for time-varying confounding factors.

Perturbation	Result
100	100
101	101
102	124
103	124
104	104
105	104
106	104
107	107
108	124
109	109
110	110
111	111
112	112
113	124
114	114
115	115
116	124
117	124
118	124
119	119
120	120
121	124
122	122
123	124
124	124
125	124
126	124
127	127
128	124
129	124
130	130
131	124
132	124
133	124
134	124
135	124 + 11 = 135
136	124
137	124
138	124
139	124

Table 17: Perturbing the first answer from Figure 5, and observing the result shows that the first answer strongly influences the final answer